Statistical Methods for High Dimensional Biology

Two group comparisons

Keegan Korthauer

25 January 2021

with slide contributions from Jenny Bryan and Sara Mostafavi

Announcements

- Seminar 2a and 2b due **today**

- Intro assignment due **today**

- Paper critique and Analysis Assignment updated

- **Initial project proposals due Thursday**

Resources for review

Reminder that (free online) resources for review of statistical concepts are listed in the syllabus

Central dogma of statistics

Image source: Josh Akey's Lecture notes

We want to understand a **population** (e.g., gene behaviour) but we can only study a **random sample** from it.

Hypothesis Testing in Genomics

Targeting of GFP to newborn rods by Nrl promoter and temporal expression profiling of flow-sorted photoreceptors

Masayuki Akimoto[*†], Hong Cheng[‡], Dongxiao Zhu[§¶], Joseph A. Brzezinski[||], Ritu Khanna[*], Elena Filippova[*], Edwin C. T. Oh[‡], Yuezhou Jing[¶], Jose-Luis Linares[*], Matthew Brooks[*], Sepideh Zareparsi[*], Alan J. Mears[*·**], Alfred Hero[§¶††‡‡], Tom Glaser[||§§], and Anand Swaroop[*‡||¶¶]

Akimoto et al. (2006)

- Retina presents a model system for investigating **regulatory networks** underlying neuronal differentiation.

- **Nrl** transcription factor is known to be important for Rod development

- **What happens if you delete Nrl?**

Why a Hypothesis Test?

From the Akimoto et al. (2006) paper:

> "we hypothesized that Nrl is the ideal transcription factor to gain insights into gene expression changes ..."

Biological question: Is the expression level of gene A affected by knockout of the *Nrl* gene?

We can use statistical inference to answer this biological question!

Statistical inference

Statistical inference:

We observe and study a random sample to make conclusions about a population (e.g., random sample of gene expressions from mice)

Experimental design:

- 4 developmental stages (E16, P2, P6, P10, 4Weeks)
- 2 genotypes: Wild type (WT), Nrl Knockout (NrlKO)
- 3-4 replicates for each combination

Reading in / exploring the data

- Data obtained from the Gene Expression Omnibus (GEO) repository

- This is explored in Seminars 4 and 5 (Seminar 5 uses the exact same data set!)

- Review lecture 3 (exploratory data analysis) for general principles

Let's take a look at 2 genes as an example: **Irs4 and Nrl**

> **Biological question:** Are these genes truly different in NrlKO compared to WT?

We can't answer this question in general. We can *only* study these genes in collected data:

We only observe a random sample of gene expression values

```
twoGenes
```

```
## # A tibble: 78 x 5
##    gene  sample_id Expression dev_stage genotype
##    <chr> <chr>          <dbl> <fct>     <fct>
##  1 Irs4  GSM92610        7.71 4_weeks   NrlKO
##  2 Irs4  GSM92611        7.77 4_weeks   NrlKO
##  3 Irs4  GSM92612        7.73 4_weeks   NrlKO
##  4 Irs4  GSM92613        7.57 4_weeks   NrlKO
##  5 Irs4  GSM92614        7.95 E16       NrlKO
##  6 Irs4  GSM92615        7.52 E16       NrlKO
##  7 Irs4  GSM92616        8.08 E16       NrlKO
##  8 Irs4  GSM92617        7.71 P10       NrlKO
##  9 Irs4  GSM92618        7.87 P10       NrlKO
```

Visualizing Irs4 and Nrl genes in our sample

Statistical Hypothesis

Experimental design: (ignoring developmental time for now)

- 2 conditions: WT *vs* NrlKO
- random sample: we observe the expression of many genes in all mice

Biological hypothesis: for *some* genes, the expression levels are different between conditions

Statistical hypotheses: (for each gene $g = 1, \ldots, G$)

- H_0 (null hypothesis): the expression level of gene *g* is the *same* in both conditions
- H_A (alternative hypothesis): the expression level of gene *g* is *different* between conditions

Notation

Random variables and estimates (we can observe):

Y_i : expression of gene g in the WT sample i

Z_i: expression of gene g in NrlKO sample i

$Y_1, Y_2, \ldots, Y_{n_Y}$: a random sample of size n_Y WT mice

$Z_1, Z_2, \ldots, Z_{n_Z}$: a random sample of size n_Z NrlKO mice

$\bar{Y} = \frac{\sum_{i=1}^{n_Y} Y_i}{n_Y}$: sample mean of gene g expression from WT mice

$\bar{Z} = \frac{\sum_{i=1}^{n_Z} Z_i}{n_Z}$: sample mean of gene g expression from NrlKO mice

Notation

Population parameters (unknown/unobservable):

$\mu_Y = E[Y]$: the (population) expected expression of gene g in WT mice

$\mu_Z = E[Z]$: the (population) expected expression of gene g in NrlKO mice

Is there **enough** evidence in the data to reject H₀?

$$H_0 : \mu_Y = \mu_Z$$

Statistical Inference: random samples are used to learn about the population

What we observe: the difference between the **sample averages**: $\bar{Y}$ vs $\bar{Z}$

```
meanExp <- twoGenes %>%
  group_by(gene, genotype) %>%
  summarize(meanExpr = mean(Expression))
meanExp
```

```
## # A tibble: 4 x 3
## # Groups:   gene [2]
##   gene  genotype meanExpr
##   <chr> <fct>       <dbl>
## 1 Irs4  NrlKO        7.74
## 2 Irs4  WT           7.77
## 3 Nrl   NrlKO        6.09
## 4 Nrl   WT          11.2
```

```
diffExp <- meanExp %>%
  group_by(gene) %>%
  summarize(diffExp = meanExpr[genotype =
                      meanExpr[genotype =
diffExp
```

```
## # A tibble: 2 x 2
##   gene  diffExp
## * <chr>   <dbl>
## 1 Irs4  -0.0261
## 2 Nrl   -5.15
```

Is the difference between $\bar{Y}$ and $\bar{Z}$ informative to reject H$_0$?
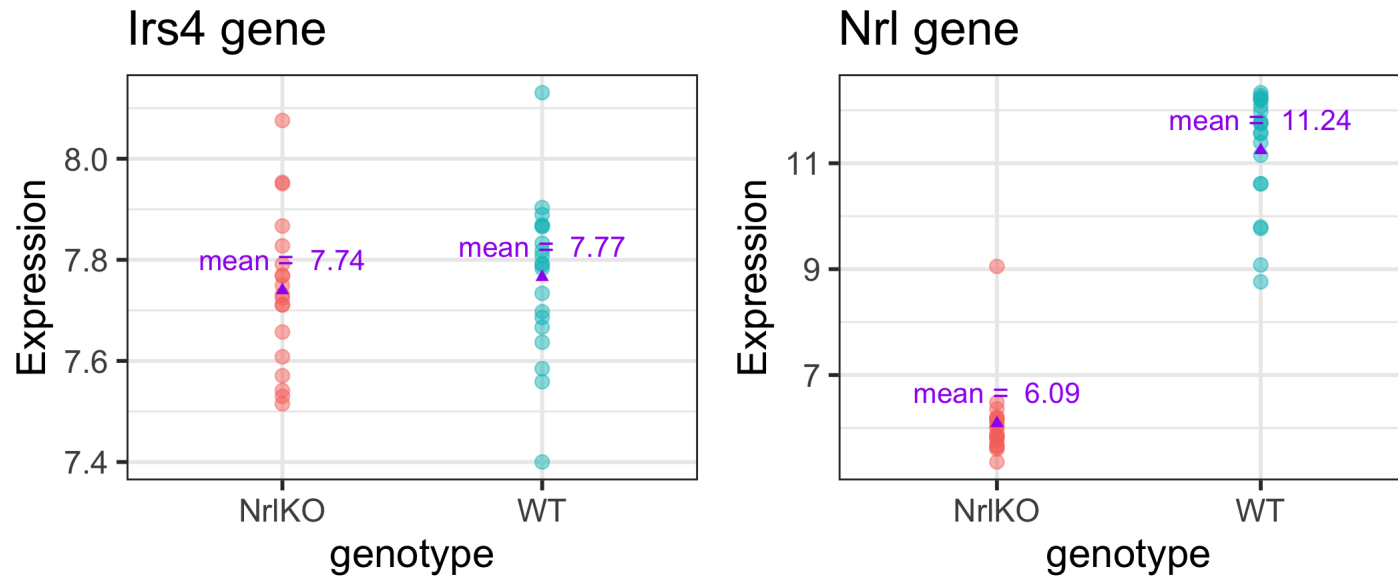
- The sample means, $\bar{Y}$ vs $\bar{Z}$, by themselves are not enough to make conclusions about the population

- What is a "large" difference? "large" relative to what?

Artificial gene

Nrl gene

What can we use to interpret the size of the mean difference?

$$\frac{\bar{Y} - \bar{Z}}{??}$$

What can we use to interpret the size of the mean difference?

"large" relative to the observed variation

$$\frac{\bar{Y}-\bar{Z}}{\sqrt{Var(\bar{Y}-\bar{Z})}}$$

Quantifying observed variation (in the difference)

- Recall that if $Var(Y_i) = \sigma_Y^2$, then $Var(\bar{Y}) = \frac{\sigma_Y^2}{n_Y}$

- Assume that the random variables within each group are *independent and identically distributed* (iid), and that the groups are independent. More specifically, that

  1. $Y_1, Y_2, \ldots, Y_{n_Y}$ are iid,
  2. $Z_1, Z_2, \ldots, Z_{n_Z}$ are iid, and
  3. Y_i, Z_j are independent. Then, it follows that

$$Var(\bar{Z} - \bar{Y}) = \frac{\sigma_Z^2}{n_Z} + \frac{\sigma_Y^2}{n_Y}$$

- If we also assume equal population variances: $\sigma_Z^2 = \sigma_Y^2 = \sigma^2$, then

$$Var(\bar{Z} - \bar{Y}) = \frac{\sigma_Z^2}{n_Z} + \frac{\sigma_Y^2}{n_Y} = \sigma^2 \left[\frac{1}{n_Z} + \frac{1}{n_Y} \right]$$

But how can we calculate population variance σ if it is **unknown?**

...using the sample variances (combined, somehow)!

```
twoGenes %>%
  group_by(gene, genotype) %>%
  summarize(groupVar = var(Expression))
```

```
## # A tibble: 4 x 3
## # Groups:   gene [2]
##   gene   genotype groupVar
##   <chr>  <fct>       <dbl>
## 1 Irs4   NrlKO      0.0233
## 2 Irs4   WT         0.0240
## 3 Nrl    NrlKO      0.594
## 4 Nrl    WT         1.22
```
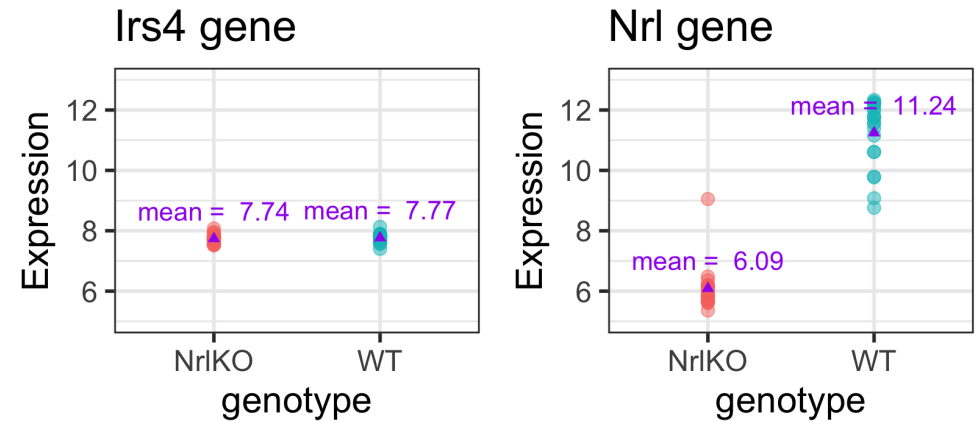

e.g., for Nrl: $\hat{\sigma}_Y^2 = S_Y^2 = \dfrac{1}{n_Y}\sum_{i=1}^{n_Y}(Y_i - \bar{Y})^2 = 1.22$

Plug these sample variances into your chosen formula for the variance of the difference of sample means

Assuming **equal** variance of Y's and Z's

$$\hat{V}(\bar{Z}_n - \bar{Y}_n) = \hat{\sigma}^2_{\text{pooled}} \left[\frac{1}{n_Y} + \frac{1}{n_Z} \right]$$

$$\hat{\sigma}^2_{\text{pooled}} = S^2_Y \frac{n_Y - 1}{n_Y + n_Z - 2} + S^2_Z \frac{n_z - 1}{n_Y + n_Z - 2}$$

Assuming **unequal** variance of Y's and Z's (Welch's t-test)

$$\hat{V}(\bar{Z}_n - \bar{Y}_n) = \hat{\sigma}^2_{\bar{Z}_n - \bar{Y}_n} = \frac{S^2_Y}{n_Y} + \frac{S^2_Z}{n_Z}$$

> Recall: the 'hat' (^) is used to distinguish an 'estimate' from a 'parameter'

The Test Statistic: $T = \dfrac{\bar{Z}_n - \bar{Y}_n}{\hat{V}(\bar{Z}_n - \bar{Y}_n)}$

```
tTests <- diffExp %>%
    mutate(t = diffExp / sqrt(s2Diff)) %>%
    mutate(tWelch = diffExp / sqrt(s2DiffWelch))
tTests
```

```
## # A tibble: 2 x 4
##    gene  diffExp        t   tWelch
## * <chr>    <dbl>    <dbl>    <dbl>
## 1 Irs4  -0.0261   -0.529   -0.529
## 2 Nrl     -5.15    -16.8    -17.0
```

Can we now say that the observed differences are 'big'?

The difference is about half a standard deviation for Irs4 and ~17 standard deviations for Nrl.

The test statistic T is a random variable because it's based on our random sample.

We need a measure of its **uncertainty** to determine how big T is:

> If we were to repeat the experiment many times, what's the probability of observing a value of T **as extreme** as the one we observed?

We need to have a probability distribution!

However, this is unknown to us so we need to **make more assumptions**.

Theory now tells us specific **null distributions** for these test statistics, depending on your assumptions.

Let's call the unknown probability distributions F and G $(Y_i \sim F, \text{ and } Z_i \sim G)$

$\Longrightarrow$ Willing to assume that F and G are normal distributions?

2-sample *t*-test:

(equal variances)

$$T \sim t_{n_Y + n_Z - 2}$$

Welch's 2-sample *t*-test:

(unequal variances)

$$T \sim t_{<something\ ugly>}$$

$\Longrightarrow$ Unwilling to assume that F and G are normal distributions? But you feel that n_Y and n_Z are large enough?

Student's *t*-distribution

Recall that T is a **random variable**. Under certain assumptions, we can prove that T follows a *t*-distribution.

where df = degrees of freedom.

Hypothesis testing: Step 1

1. Formulate your hypothesis as a statistical hypothesis:

$$H_0 : \mu_Y = \mu_Z \text{ vs } H_A : \mu_Y \neq \mu_Z$$

Hypothesis testing: Step 2

2a. Define a **test statistic**: 2-sample *t*-test

2b. Compute the observed value for the test statistic:

```
twoGenes %>%
  group_by(gene) %>%
  summarize(t = t.test(Expression ~ genotype,
                       var.equal=TRUE)$statistic)
```

```
## # A tibble: 2 x 2
##   gene        t
## * <chr>   <dbl>
## 1 Irs4   -0.529
## 2 Nrl    -16.8
```

Hypothesis testing: Step 3

3. Compute the probability of seeing a test statistic at least as extreme as that observed, under the **null sampling distribution** (this is the definition of the p-value)

```
twoGenes %>%
  group_by(gene) %>%
  summarize(pvalue = t.test(Expression ~ genotype,
                            var.equal=TRUE)$p.value)
```

```
## # A tibble: 2 x 2
##   gene    pvalue
## * <chr>    <dbl>
## 1 Irs4  6.00e- 1
## 2 Nrl   6.73e-19
```

In other words, assuming that H_0 is true:

For Irs4, the probability of seeing a test statistic as extreme as that observed $(t = -0.53)$ is pretty high $(p = 0.6)$.

But for Nrl, the probability of seeing a test statistic as extreme as that observed $(t = -16.8)$ is extremely low $(p = 6.76 \times 10^{-19})$
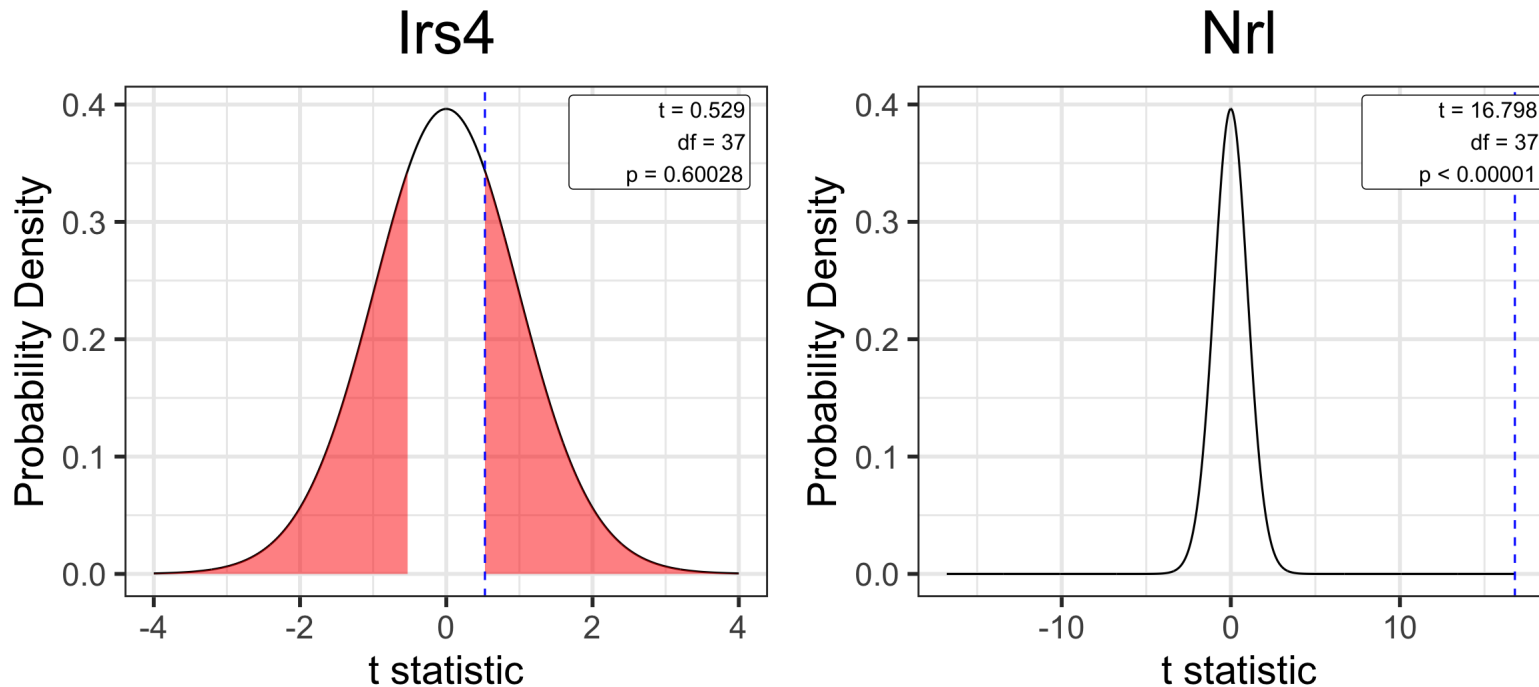
Hypothesis Testing: Step 4

4. Make a decision about significance of results, based on a pre-specified value (alpha, significance level)

The significance level α is often set at 0.05. However, this value is arbitrary and may depend on the study.

Irs4

Using $\alpha = 0.05$, since the p-value for the Irs4 test is greater than 0.05, we conclude that there is not enough evidence in the data to claim that Irs4 has differential expression in WT compared to NrlKO models.

We do not reject H_0!

Nrl

Using $\alpha = 0.05$, since the p-value for the Nrl test is much less than 0.05, we conclude that there is significant evidence in the data to claim that Nrl has differential expression in WT compared to NrlKO models.

We reject H_0!

`t.test` function in R

Assuming equal variances

```
t.test(Expression ~ genotype,
       var.equal=TRUE,
       data = twoGenes  %>% filter(gene == "Nrl"))
```

```
##
## Two Sample t-test
##
## data: Expression by genotype
## t = -16.798, df = 37, p-value < 2.2e-16
## alternative hypothesis: true difference in
means is not equal to 0
## 95 percent confidence interval:
## -5.776672 -4.533071
## sample estimates:
## mean in group NrlKO mean in group WT
## 6.089579 11.244451
```

Assuming equal variances

```
t.test(Expression ~ genotype,
       var.equal=FALSE,
       data = twoGenes  %>% filter(gene == "Nrl"))
```

```
##
## Welch Two Sample t-test
##
## data: Expression by genotype
## t = -16.951, df = 34.01, p-value < 2.2e-16
## alternative hypothesis: true difference in
means is not equal to 0
## 95 percent confidence interval:
## -5.772864 -4.536879
## sample estimates:
## mean in group NrlKO mean in group WT
## 6.089579 11.244451
```

What is a p-value?

> Likelihood of obtaining a test statistic at least as extreme as the one observed, given that the null hypothesis is true (we are making a *conditional probability* statement)

What is a p-value NOT?

- Not the probability that the null hypothesis is true

- Not the probability that the finding is a "fluke"

- Not the probability of falsely rejecting the null

- Does not indicate the size or importance of observed effects.

"Genome-wide" testing of differential expression

- In genomics, we often perform thousands of statistical tests (e.g., a *t*-test per gene)

- The distribution of p-values across all tests provides good diagnostics/insights.

- Is it uniform (should be in most experiments)? If not, is the departure from uniform expected based on biological knowledge?

- We will come back to this in greater detail in a later lecture

Different kinds of *t*-tests:

- One sample *or* **two samples**

- One-sided *or* **two sided**

- Paired *or* **unpaired**

- Equal variance *or* unequal variance

Types of Errors in Hypothesis Testing

Actual Situation "Truth"

Decision	H_0 True	H_0 False
Do Not Reject H_0	Correct Decision $1-\alpha$	Incorrect Decision Type II Error β
Reject H_0	Incorrect Decision Type I Error α	Correct Decision $1-\beta$

$$\alpha = P(\text{Type I Error}), \ \beta = P(\text{Type II Error}), \ \text{Power} = 1 - \beta$$

H_0: "Innocent until proven guilty"

- The default state is $H_0 \rightarrow$ we only reject if we have enough evidence

- If H_0: Innocent and H_A: Guilty, then

  - Type I Error (α): Wrongfully convict innocent (*False Positive*)

  - Type II Error (β): Fail to convict criminal (*False Negative*)

What are alternatives to the *t*-test?

> What if you don't wish to assume the underlying data is normally distributed **AND** you aren't sure your samples are large enough to invoke CLT?

First, one could use the t test statistic but use a **bootstrap approach** to compute its p-value. We will cover this later on.

Alternatively, there are *non-parametric* tests that are available here:

- **Wilcoxon rank sum test**, aka Mann Whitney, uses ranks to test differences in population means.

- **Kolmogorov-Smirnov test** uses the empirical CDF to test differences in population cumulative distributions.

Wilcoxon rank sum test

Rank all data, ignoring the **grouping** variable

Test statistic = sum of the ranks for one group (optionally, subtract the minimum possible which is $\frac{n_Y(n_Y+1)}{2}$)

(Alternative but equivalent formulation based on the number of y_i, z_i pairs for which $y_i \geq z_i$)

Null distribution of such statistics can be worked out or approximated.

`wilcox.test` function in R

```
wilcox.test(Expression ~ genotype,
            data = twoGenes  %>% filter(gene == "Irs4"))
```

```
##
##      Wilcoxon rank sum exact test
##
## data:  Expression by genotype
## W = 160, p-value = 0.4115
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(Expression ~ genotype,
            data = twoGenes  %>% filter(gene == "Nrl"))
```

```
##
##      Wilcoxon rank sum exact test
##
## data:  Expression by genotype
## W = 1, p-value = 5.804e-11
## alternative hypothesis: true location shift is not equal to 0
```

Kolmogorov-Smirnov test (two sample)

Null hypothesis: F = G, i.e. the distributions are the same

Estimate each CDF with the empirical CDF (ECDF)

$$\hat{F}(x) = \frac{1}{n} \sum_{i=1}^{n} I[x_i \leq x]$$

Test statistic is the maximum of the absolute difference between the ECDFs
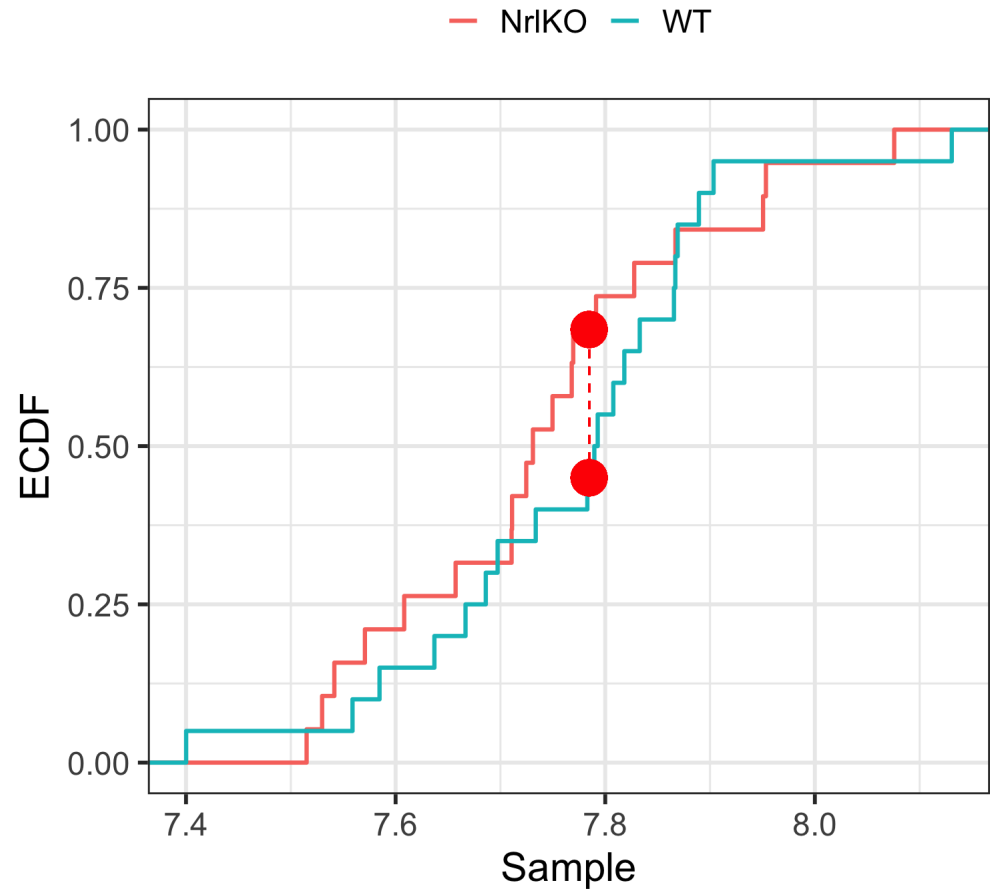
$$max|\hat{F}(x) - \hat{G}(x)|$$

Null distribution does not depend on F, G (!)
(I'm suppressing detail here)

Kolmogorov-Smirnov test (two sample)

`ks.test` function in R

```
Nrlgene <- twoGenes %>% filter(gene == "Nrl")
ks.test(Nrlgene$Expression[Nrlgene$genotype == "WT"],
        Nrlgene$Expression[Nrlgene$genotype == "NrlKO"])
```

```
##
## Two-sample Kolmogorov-Smirnov test
##
## data: Nrlgene$Expression[Nrlgene$genotype == "WT"] and
Nrlgene$Expression[Nrlgene$genotype == "NrlKO"]
## D = 0.95, p-value = 5.804e-10
## alternative hypothesis: two-sided
```

```
Irs4gene <- twoGenes %>% filter(gene == "Irs4")
ks.test(Irs4gene$Expression[Irs4gene$genotype == "WT"],
        Irs4gene$Expression[Irs4gene$genotype == "NrlKO"])
```

```
##
## Two-sample Kolmogorov-Smirnov test
##
## data: Irs4gene$Expression[Irs4gene$genotype == "WT"] and
Irs4gene$Expression[Irs4gene$genotype == "NrlKO"]
## D = 0.28421, p-value = 0.3278
## alternative hypothesis: two-sided
```

Discussion and questions ...

- What if you are unsure whether your sample size is large enough? Outliers with small samples could be problematic

- Which test result should one report ... the 2-sample t-test, the Wilcoxon, or the KS?

- Treat p-values as one type of evidence that you should incorporate with others

- It is worrisome when methods that are equally appropriate and defensible give very different answers