Statistical Methods for High Dimensional Biology

## More limma & multiple testing

Keegan Korthauer

8 February 2021

with slide contributions from Gabriela Cohen Freue, Jenny Bryan, and Sara Mostafavi

Announcements

- Feedback posted on your initial project proposals (see the issue in your team's repository)

- Reminder that the finalized project proposals are due Feb 18

- Instructions on how to submit your deliverables have been updated for clarity on the website

  - in particular pay attention to the recommended way to grab the link to the submitted files from your latest commit

Course format feedback
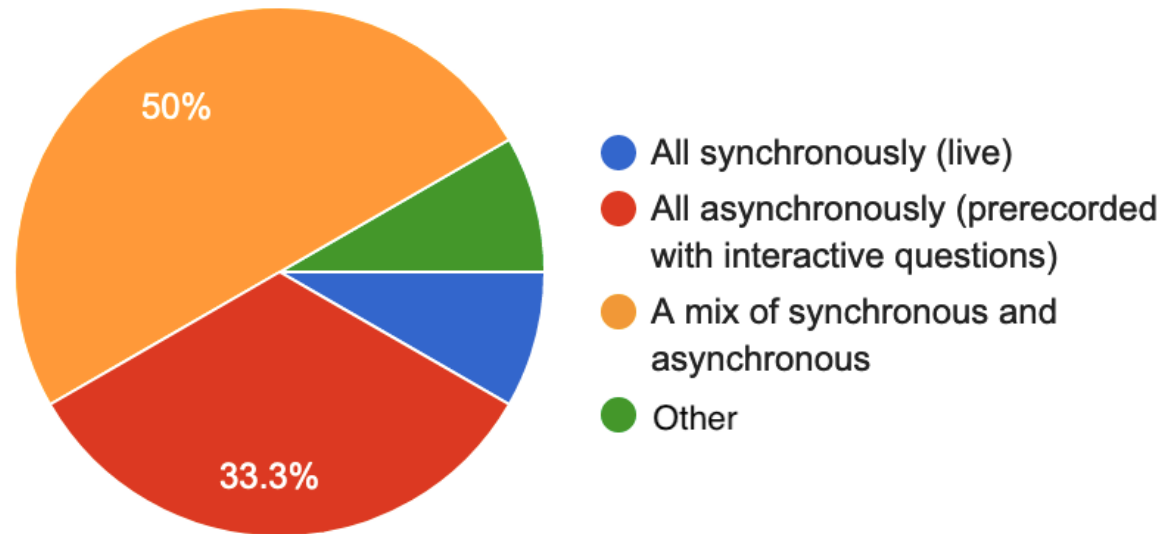
- Preference for mix of synchronous & asynchronous

- 😊 Positive feedback on interactive quiz questions / polls / chat

- 🙁 Several people felt a lack of communication - let's utilize the Discussion repo more!

Learning objectives

- How to use `limma` to perform genome-wide differential expression testing on microarray data

- Understand the key differences between `limma` and `lm`

- Explain why multiple testing increases the number of errors we make by chance

- Be able to adjust for multiple comparisons by controlling the False Discovery Rate

  - e.g. using Benjamini-Hochberg or Storey's q-value

Recall: The hybrid estimator in `limma`

$$\tilde{s}_g^2 = \frac{d_0 s_0^2 + d s_g^2}{d_0 + d}$$

- recall that (s_0, d_0) are the *prior* parameters for σ_g^2 (random variable):

$$\frac{1}{\sigma_g^2} \sim \frac{1}{d_0 s_0^2} \chi_{d_0}^2$$

- the prior parameters incorporate information from all genes which allows us to **shrink/nudge the gene-specific variances toward a common consensus**

  - they are estimated from the data - the formulas for (s_0, d_0) and their derivations are beyond the scope of the course (but `limma` takes care of the details for us)

  - note that (s_0, d_0) do not depend on g

`limma` quickstart

$$\mathbf{Y}_g = \mathbf{X}\boldsymbol{\alpha}_g + \boldsymbol{\varepsilon}_g,$$

$$E(\boldsymbol{\varepsilon}_g) = 0, \ Var(\boldsymbol{\varepsilon}_g) = \sigma_g^2, \ \varepsilon_{ig} \perp \varepsilon_{jg},$$

- Within each gene observations are iid / constant variance

- `lmFit()` function in `limma` carries out multiple linear regression on each gene

- Usage: `lmFit(myDat, desMat)`

  - `myDat` is a data frame, matrix with features in rows and samples in columns (G genes by N samples), or `ExpressionSet` object

  - `desMat` is a design matrix (output of `model.matrix(y ~ x)`; N samples by p parameters)
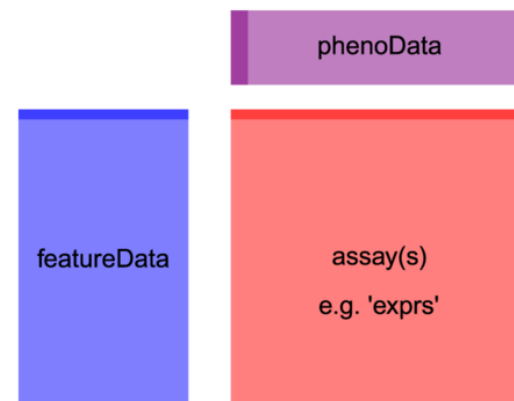

Bioconductor's ExpressionSet class

Let's run `limma` for the interactive model with age

$$y_{ig} = \theta + \tau_{KO} x_{ig,KO} + \tau_{Age} x_{ig,Age} + \tau_{KO:Age} x_{ig,KO} x_{ig,Age}$$

- i indexes mouse, g indexes genes

- $x_{ig,KO}$ is the dummy/indicator variable for the NrlKO group

- $x_{ig,Age}$ is the continuous age variable

Let's run `limma` for the interactive model with age

Example gene (but we want to fit this model on all genes):

Tmem176a

Arranging input for `lmFit`: Bioconductor way

```
eset
```

```
## ExpressionSet (storageMode: lockedEnvironment)
## assayData: 45101 features, 39 samples
##   element names: exprs
## protocolData: none
## phenoData
##   sampleNames: GSM92610 GSM92611 ... GSM92648 (39 total)
##   varLabels: sample_id dev_stage genotype age
##   varMetadata: labelDescription
## featureData: none
## experimentData: use 'experimentData(object)'
##   pubMedIds: 16505381
## Annotation: GPL1261
```

Arranging input for `lmFit`: Separate expression & metadata

```
myDat <- assayData(eset)$exprs
myDat[1:5,1:5]
```

```
##             GSM92610 GSM92611  GSM92612 GSM92613  GSM92614
## 1415670_at  7.108863 7.322392  7.420947 7.351444  7.240428
## 1415671_at  9.714002 9.797742  9.831072 9.658442  9.708906
## 1415672_at  9.429030 9.846977 10.003092 9.914112 10.174548
## 1415673_at  8.426974 8.404206  8.594600 8.404206  8.835334
## 1415674_a_at 8.498338 8.458287  8.426651 8.372776  8.541722
```

```
myMeta <- pData(eset)
myMeta[1:5,]
```

```
##           sample_id dev_stage genotype age
## GSM92610  GSM92610         4W    NrlKO  28
## GSM92611  GSM92611         4W    NrlKO  28
## GSM92612  GSM92612         4W    NrlKO  28
## GSM92613  GSM92613         4W    NrlKO  28
## GSM92614  GSM92614        E16    NrlKO  -4
```

Formulating input for `lmFit`: Design Matrix

Bioconductor way:

```
desMat <- model.matrix(~ genotype*age,
                       data = pData(eset))
```

Equivalently, if using the separate way:

```
desMat <- model.matrix(~ genotype*age,
                       data = myMeta)
```

`desMat`

```
##           (Intercept) genotypeNrlKO age genotypeNrlKO:age
## GSM92610            1             1  28                28
## GSM92611            1             1  28                28
## GSM92612            1             1  28                28
## GSM92613            1             1  28                28
## GSM92614            1             1  -4                -4
## GSM92615            1             1  -4                -4
## GSM92616            1             1  -4                -4
## GSM92617            1             1  10                10
## GSM92618            1             1  10                10
## GSM92619            1             1  10                10
## GSM92620            1             1  10                10
## GSM92621            1             1   2                 2
## GSM92622            1             1   2                 2
## GSM92623            1             1   2                 2
## GSM92624            1             1   2                 2
## GSM92625            1             1   6                 6
## GSM92626            1             1   6                 6
## GSM92627            1             1   6                 6
## GSM92628            1             1   6                 6
## GSM92629            1             0  28                 0
## GSM92630            1             0  28                 0
## GSM92631            1             0  28                 0
## GSM92632            1             0  28                 0
## GSM92633            1             0  -4                 0
## GSM92634            1             0  -4                 0
## GSM92635            1             0  -4                 0
## GSM92636            1             0  -4                 0
## GSM92637            1             0  10                 0
## GSM92638            1             0  10                 0
```

Computation is fast

```
system.time(gFit <- lmFit(eset,
                          model.matrix(~ genotype*age, data = pData(eset))))
```

```
##    user  system elapsed
##   0.108   0.023   0.133
```

- Equivalently, using the 'separate' way:

  ```
  gFit <- lmFit(myDat,
              model.matrix(~ genotype*age, data = myMeta))
  ```

- Using `lmFit` to fit an interactive model on 45K probesets takes a fraction of a second

- The time-intensive parts of an analysis lie in selecting the model and covariates, choosing how to parameterize it, and interpreting the output

Output of `lmFit`

```
summary(gFit)
```

```
##                      Length Class  Mode
## coefficients        180404 -none- numeric
## rank                     1 -none- numeric
## assign                   4 -none- numeric
## qr                       5 qr     list
## df.residual          45101 -none- numeric
## sigma                45101 -none- numeric
## cov.coefficients        16 -none- numeric
## stdev.unscaled      180404 -none- numeric
## pivot                    4 -none- numeric
## Amean                45101 -none- numeric
## method                   1 -none- character
## design                 156 -none- numeric
```

```
nrow(eset)
```

```
## Features
##    45101
```

```
nrow(eset)*4
```

```
## Features
##   180404
```

Output of `lmFit`

```
summary(gFit)
```

```
##                     Length Class  Mode
## coefficients        180404 -none- numeric
## rank                     1 -none- numeric
## assign                   4 -none- numeric
## qr                       5 qr     list
## df.residual          45101 -none- numeric
## sigma                45101 -none- numeric
## cov.coefficients        16 -none- numeric
## stdev.unscaled      180404 -none- numeric
## pivot                    4 -none- numeric
## Amean                45101 -none- numeric
## method                   1 -none- character
## design                 156 -none- numeric
```

```
nrow(eset)
```

```
## Features
##    45101
```

```
nrow(eset)*4
```

```
## Features
##   180404
```

- OK... but where are the shrunken variable estimates?? How do I pull out p-values??

- Actually, we haven't carried out the empirical Bayesian computation yet -- need to run `eBayes()`!

eBayes()

summary(gFit)

```
##                  Length Class  Mode
## coefficients     180404 -none- numeric
## rank                  1 -none- numeric
## assign                4 -none- numeric
## qr                    5 qr     list
## df.residual       45101 -none- numeric
## sigma             45101 -none- numeric
## cov.coefficients     16 -none- numeric
## stdev.unscaled   180404 -none- numeric
## pivot                 4 -none- numeric
## Amean             45101 -none- numeric
## method                1 -none- character
## design              156 -none- numeric
```

summary(ebFit <- eBayes(gFit))

```
##                  Length Class  Mode
## coefficients     180404 -none- numeric
## rank                  1 -none- numeric
## assign                4 -none- numeric
## qr                    5 qr     list
## df.residual       45101 -none- numeric
## sigma             45101 -none- numeric
## cov.coefficients     16 -none- numeric
## stdev.unscaled   180404 -none- numeric
## pivot                 4 -none- numeric
## Amean             45101 -none- numeric
## method                1 -none- character
## design              156 -none- numeric
## df.prior              1 -none- numeric
## s2.prior              1 -none- numeric
## var.prior             4 -none- numeric
## proportion            1 -none- numeric
## s2.post           45101 -none- numeric
## t                180404 -none- numeric
## df.total          45101 -none- numeric
## p.value          180404 -none- numeric
## lods             180404 -none- numeric
## F                 45101 -none- numeric
## F.p.value         45101 -none- numeric
```

Components of the empirical Bayes estimators

math	plain english	`limma`	numerical result	also in `lm`?
s_g^2	gene-specific residual variance	`gFit$sigma^2`	45K numbers	✓
d	residual degrees of freedom $(n - p)$	`gFit$df.residual`	$39 - 4 = 35^*$	✓
s_0^2	mean of inverse χ^2 prior for s_g^2	`ebFit$s2.prior`	0.048	
d_0	degrees of freedom for the prior	`ebFit$df.prior`	2.61	
$\tilde{s}_g^2$	posterior mean of s_g^2 (i.e. moderated residual variance)	`ebFit$s2.post`	45K numbers	

* limma can handle more complicated models where this is not the same for each gene, so this is actually a vector of 45K copies of the number 35

`topTable()` will help us extract relevant output in a convenient format!

```
topTable(fit, coef=NULL, number=10, genelist=fit$genes, adjust.method="BH",
         sort.by="B", resort.by=NULL, p.value=1, lfc=0, confint=FALSE)
```

`fit`	list containing a linear model fit produced by `lmFit`, `lm.series`, `gls.series` or `mrlm`. For `topTable`, `fit` should be an object of class `MArrayLM` as produced by `lmFit` and `eBayes`.
`coef`	column number or column name specifying which coefficient or contrast of the linear model is of interest. For `topTable`, can also be a vector of column subscripts, in which case the gene ranking is by F-statistic for that set of contrasts.
`number`	maximum number of genes to list
`adjust.method`	method used to adjust the p-values for multiple testing. Options, in increasing conservatism, include "none", "BH", "BY" and "holm". See `p.adjust` for the complete list of options. A `NULL` value will result in the default adjustment method, which is "BH".
`sort.by`	character string specifying statistic to rank genes by. Possible values for `topTable` and `toptable` are "logFC", "AveExpr", "t", "P", "p", "B" or "none". (Permitted synonyms are "M" for "logFC", "A" or "Amean" for "AveExpr", "T" for "t" and "p" for "P".) Possibilities for `topTableF` are "F" or "none". Possibilities for `topTreat` are as for `topTable` except for "B".

… (truncated - see `?topTable` for full listing)

Summary of `topTable` function

- `coef` is the argument where you specify the coefficient(s) you want to test for equality with zero (default is NULL; must be specified)

- `number` lets you control size of hit list (default is 10)

- `p.value` lets you specify a minimum adjusted p-value cutoff (default is 1)

- `lfc` lets you specify a minimum observed effect size - log2 fold change (default is 0)

- `sort.by` and `resort.by` give control over the ordering (default is by "B": log-odds that the gene is differentially expressed)

- `adjust.method` specifies how/if to adjust p-values for multiple testing (default is BH)

`topTable` in action: `genotypeNrlKO`

```
topTable(ebFit, coef = "genotypeNrlKO")
```

```
##                    logFC  AveExpr           t       P.Value     adj.P.Val          B
## 1450946_at    -4.7134891 8.733103  -15.583683 5.390931e-18 2.431364e-13 28.40938
## 1442070_at    -0.9669400 8.268122   -9.787407 6.896220e-12 1.555132e-07 16.46086
## 1443457_at    -1.1530057 6.049428   -9.091412 4.947029e-11 6.029428e-07 14.68843
## 1422679_s_at  -2.1187394 9.495435   -9.002625 6.388954e-11 6.029428e-07 14.45703
## 1431708_a_at  -2.1610625 8.591450   -8.967950 7.062007e-11 6.029428e-07 14.36634
## 1433050_at    -1.5083887 8.468891   -8.923932 8.021234e-11 6.029428e-07 14.25096
## 1426288_at    -4.2100657 9.323796   -8.772209 1.246491e-10 8.031144e-07 13.85106
## 1418108_at     0.8662091 8.260647    8.537984 2.475294e-10 1.395478e-06 13.22713
## 1450770_at     1.3103576 7.357033    8.128847 8.332369e-10 3.747976e-06 12.11870
## 1457802_at    -0.8362030 7.778189   -8.117112 8.629984e-10 3.747976e-06 12.08657
```

- `topTable(ebFit, coef = 2)` is equivalent here, but much less informative!!

- What is the null hypothesis here?

  - $H_0 : \tau_{KO} = 0$

  - this finds genes where the knockouts differ from the wild types *when age is zero*
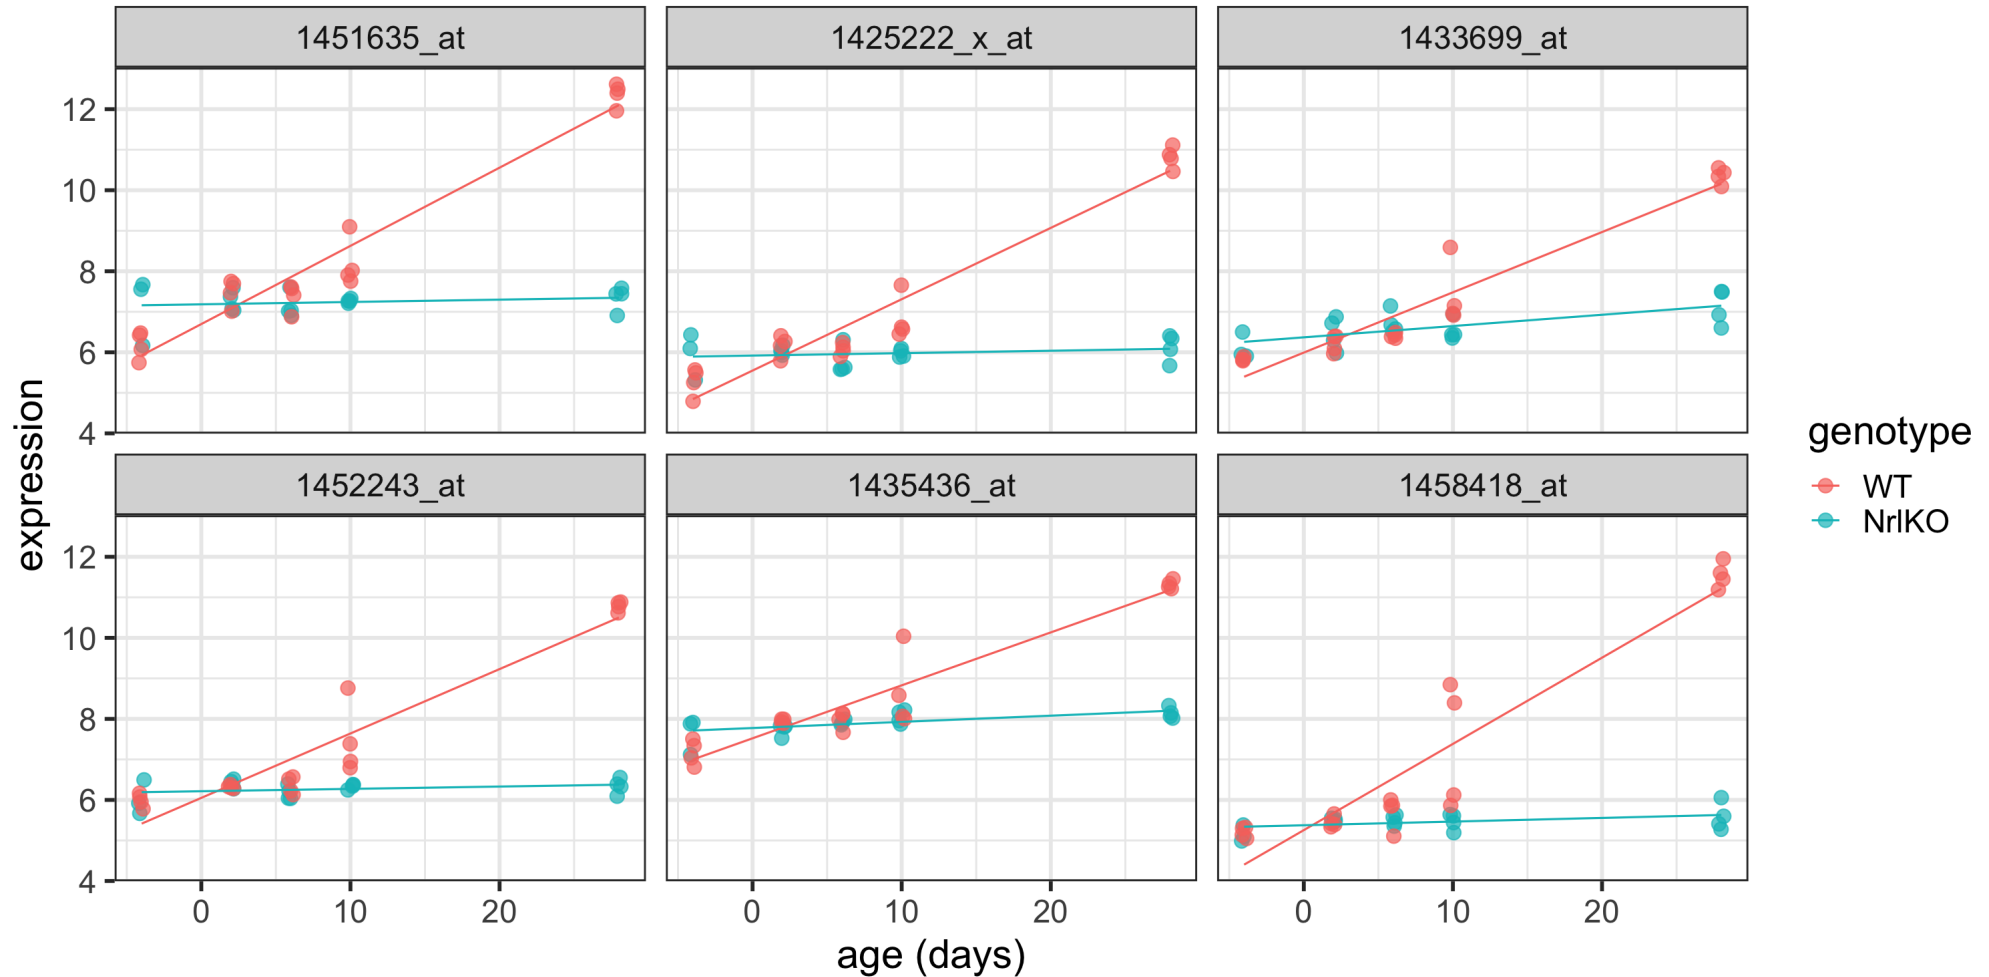
Plotting the top 6 probes for `genotypeNrlKO`

`topTable` in action: `age`

```
topTable(ebFit, coef = "age")
```

```
##                   logFC  AveExpr        t      P.Value   adj.P.Val        B
## 1451635_at   0.19294061 7.791237 20.78935 3.204717e-22 1.445359e-17 40.60768
## 1425222_x_at 0.17604111 6.514475 19.72633 1.962954e-21 4.426561e-17 38.79502
## 1433699_at   0.14875031 6.939973 17.80467 6.442530e-20 7.957605e-16 35.29412
## 1452243_at   0.15889179 6.840975 17.75658 7.057586e-20 7.957605e-16 35.20253
## 1435436_at   0.13067119 8.274676 17.61808 9.187120e-20 8.286966e-16 34.93763
## 1458418_at   0.21277027 6.270187 16.43686 9.326171e-19 7.006477e-15 32.60718
## 1424977_at   0.07814988 6.505464 16.36079 1.087456e-18 7.006477e-15 32.45260
## 1431174_at   0.16158579 7.418000 16.01901 2.183055e-18 1.230725e-14 31.75106
## 1421818_at   0.11250919 7.982278 15.66666 4.531133e-18 2.270651e-14 31.01562
## 1419069_at   0.09325234 8.222096 15.54775 5.813459e-18 2.621928e-14 30.76457
```

- `topTable(ebFit, coef = 3)` is equivalent here, but much less informative!!

- What is the null hypothesis here?

  - $H_0 : \tau_{Age} = 0$

  - this finds genes where age significantly affects gene expression *for WT*

Plotting the top 6 genes for `age`

`topTable` in action: `genotypeNrlKO:age`

```
topTable(ebFit, coef = "genotypeNrlKO:age")
```

```
##                   logFC  AveExpr         t      P.Value    adj.P.Val         B
## 1451635_at   -0.1871800 7.791237 -14.00463 1.684302e-16 7.596369e-12 27.52010
## 1425222_x_at -0.1700908 6.514475 -13.23451 9.961393e-16 2.246344e-11 25.74877
## 1441809_at    0.1336356 6.649440  12.55702 5.037221e-15 7.572790e-11 24.13110
## 1452243_at   -0.1531326 6.840975 -11.88284 2.670693e-14 3.011273e-10 22.46414
## 1431174_at   -0.1655894 7.418000 -11.39882 9.161231e-14 8.263614e-10 21.23123
## 1458418_at   -0.2037314 6.270187 -10.92850 3.122990e-13 2.347500e-09 20.00384
## 1435436_at   -0.1154500 8.274676 -10.80855 4.289363e-13 2.763637e-09 19.68613
## 1446715_at   -0.1409774 5.613122 -10.57909 7.912260e-13 4.325169e-09 19.07307
## 1442190_at   -0.1323038 5.896077 -10.54672 8.630965e-13 4.325169e-09 18.98601
## 1416306_at    0.1702059 6.560050  10.41285 1.238250e-12 5.584630e-09 18.62456
```

- `topTable(ebFit, coef = 4)` is equivalent here, but much less informative!!

- What is the null hypothesis here?

  - $H_0 : \tau_{KO:Age} = 0$

  - this finds genes where the effect of age is significantly different in each genotype

# Plotting the top 6 genes for `genotypeNrlKO:age`
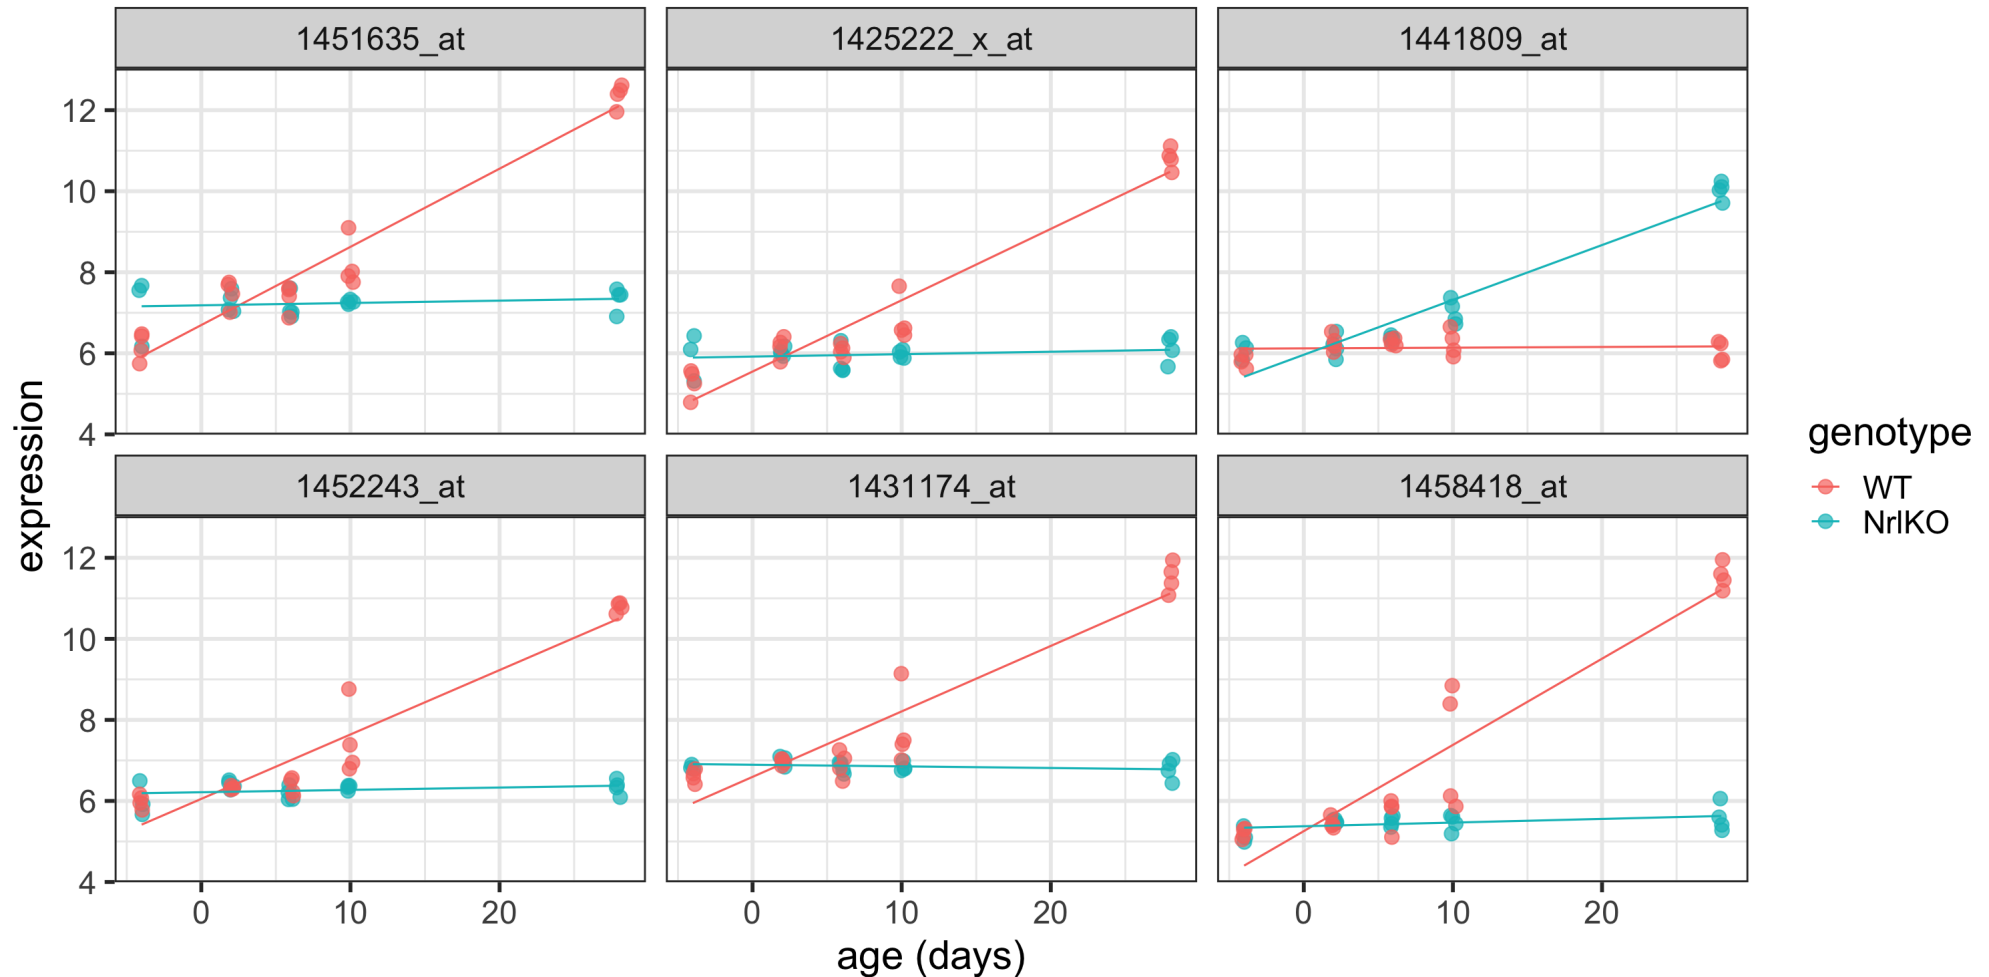
`topTable` in action: any effect of genotype

```
topTable(ebFit, coef = c("genotypeNrlKO", "genotypeNrlKO:age"))[,-3]
```

```
##                genotypeNrlKO genotypeNrlKO.age        F      P.Value    adj.P.Val
## 1450946_at       -4.7134891       -0.05411298 246.5762 2.403638e-22 1.084065e-17
## 1451635_at        0.4868705       -0.18718005 129.8112 1.307185e-17 2.180049e-13
## 1451617_at       -0.9885684       -0.22329233 128.9935 1.450111e-17 2.180049e-13
## 1425222_x_at      0.3698846       -0.17009077 120.0041 4.719613e-17 5.321482e-13
## 1441809_at       -0.1576064        0.13363563 117.5571 6.594308e-17 5.948198e-13
## 1416306_at        0.4278776        0.17020589 113.2125 1.212332e-16 9.112895e-13
## 1439143_at        0.6570269        0.06170447 111.2102 1.616023e-16 1.001158e-12
## 1449526_a_at      0.9935801        0.04232958 110.5598 1.775851e-16 1.001158e-12
## 1452243_at        0.1639373       -0.15313265 106.2784 3.344140e-16 1.675823e-12
## 1451763_at       -0.6819899       -0.14010359 102.7010 5.771358e-16 2.442992e-12
```

- `topTable(ebFit, coef = c(2,4))` is equivalent here, but much less informative!!

- What is the null hypothesis here?

  - $H_0 : \tau_{KO} = \tau_{KO:Age} = 0$

  - this finds genes where any (additive/interaction) effect of genotype is significant

Plotting the top 6 genes for any effect of genotype
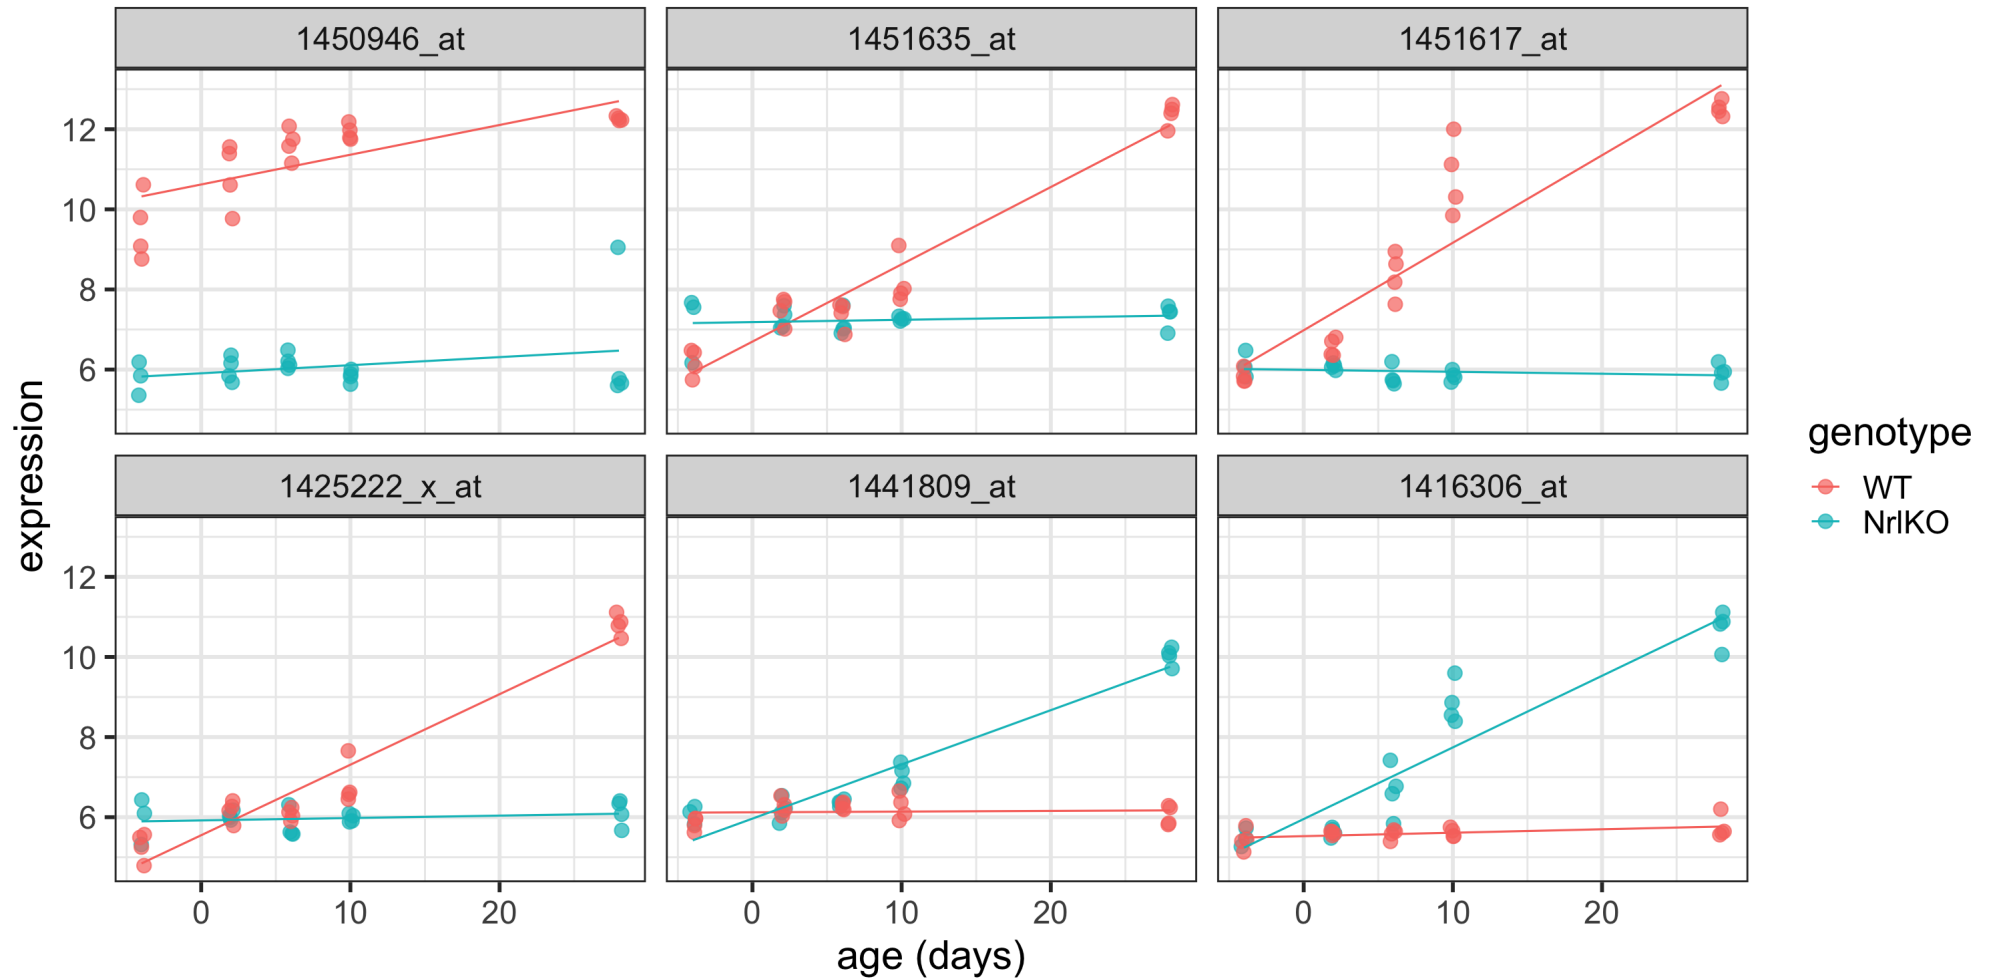
Comparison of s_g^2 and $\tilde{s}_g^2$ (shrinkage!)

For **small** variances, limma _____ the estimates (fill in the blank):

(a) INCREASES

(b) DECREASES

Ests of gene-specific variance

Comparison of s_g^2 and $\tilde{s}_g^2$ (shrinkage!)

For **large** variances, limma _____ the estimates (fill in the blank):

(a) INCREASES

(b) DECREASES

Ests of gene-specific variance

Comparison of interaction coefficient p-values
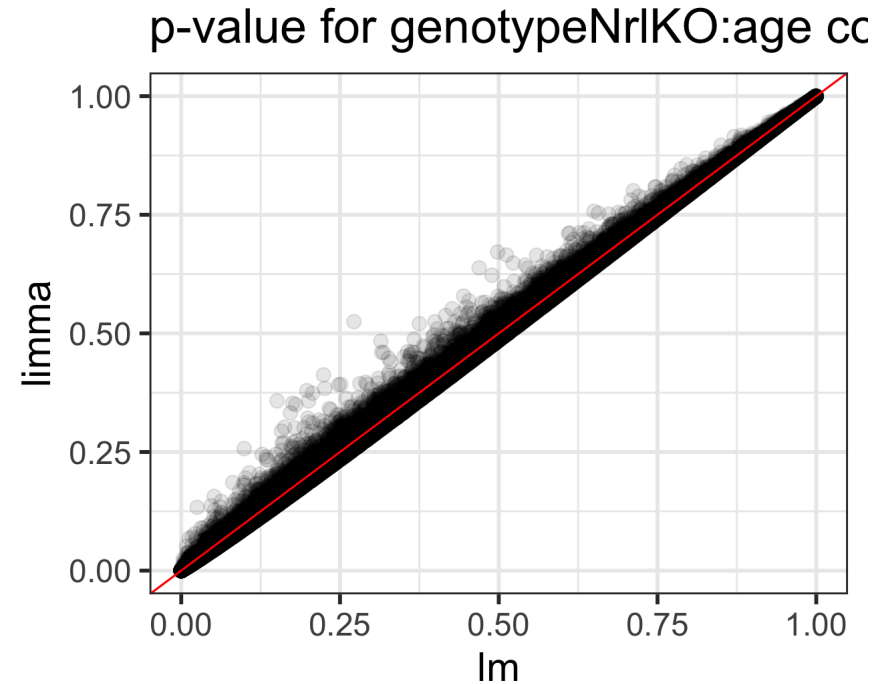
- 17479 genes where limma p-value is *larger* than lm
- 27622 genes where limma p-value is *smaller* than lm

Multiple testing

Error rates

Actual Situation "Truth"

Decision	H_0 True	H_0 False
Do Not Reject H_0	Correct Decision 1-α	Incorrect Decision Type II Error β
Reject H_0	Incorrect Decision Type I Error α	Correct Decision 1-β

$$\alpha = P(\text{Type I Error}), \ \beta = P(\text{Type II Error}), \ \text{Power} = 1 - \beta$$

Type I Error rate for m tests

- $P(\text{incorrect decision}|H_0) = \alpha$

  - let $\alpha = 0.05$

- $P(\text{correct decision}|H_0) = 1 - \alpha = 0.95$

Type I Error rate for m tests

- $P(\text{incorrect decision}|H_0) = \alpha$

    ○ let $\alpha = 0.05$

- $P(\text{correct decision}|H_0) = 1 - \alpha = 0.95$

- $P(\text{correct decision on } m \text{ tests}|H_0) =$

$$(1 - \alpha)^m = 0.95^m$$

-

    $P(\text{at least one incorrect decision on } m \text{ tests}|H_0) =$

    $$1 - (1 - \alpha)^m =$$
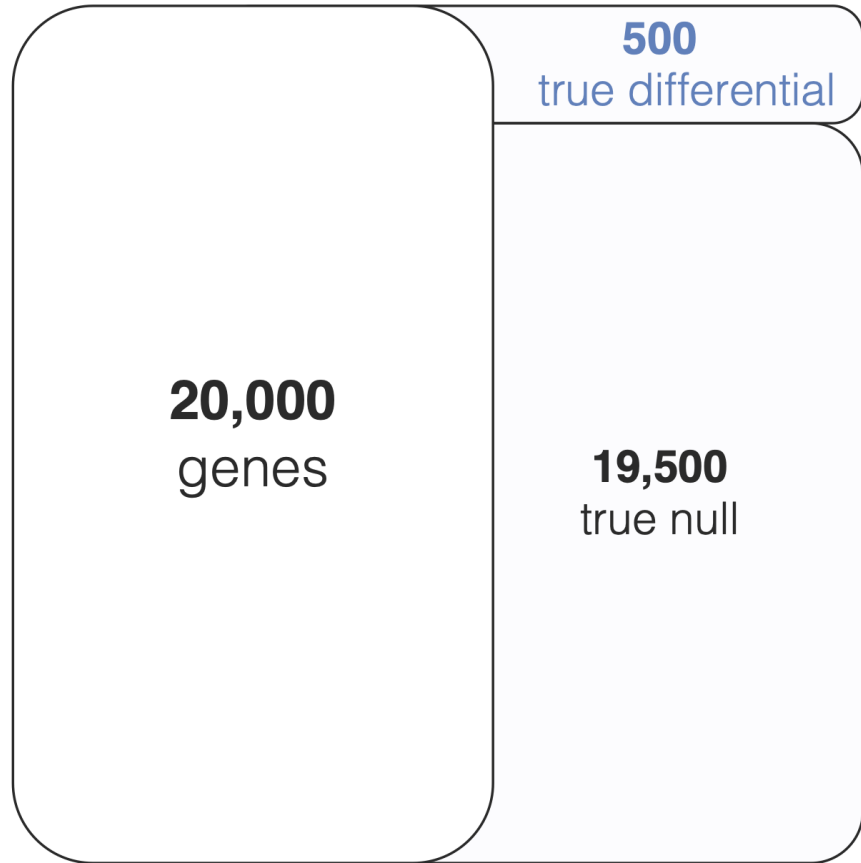
    $$1 - 0.95^m = \alpha_{FWER}$$
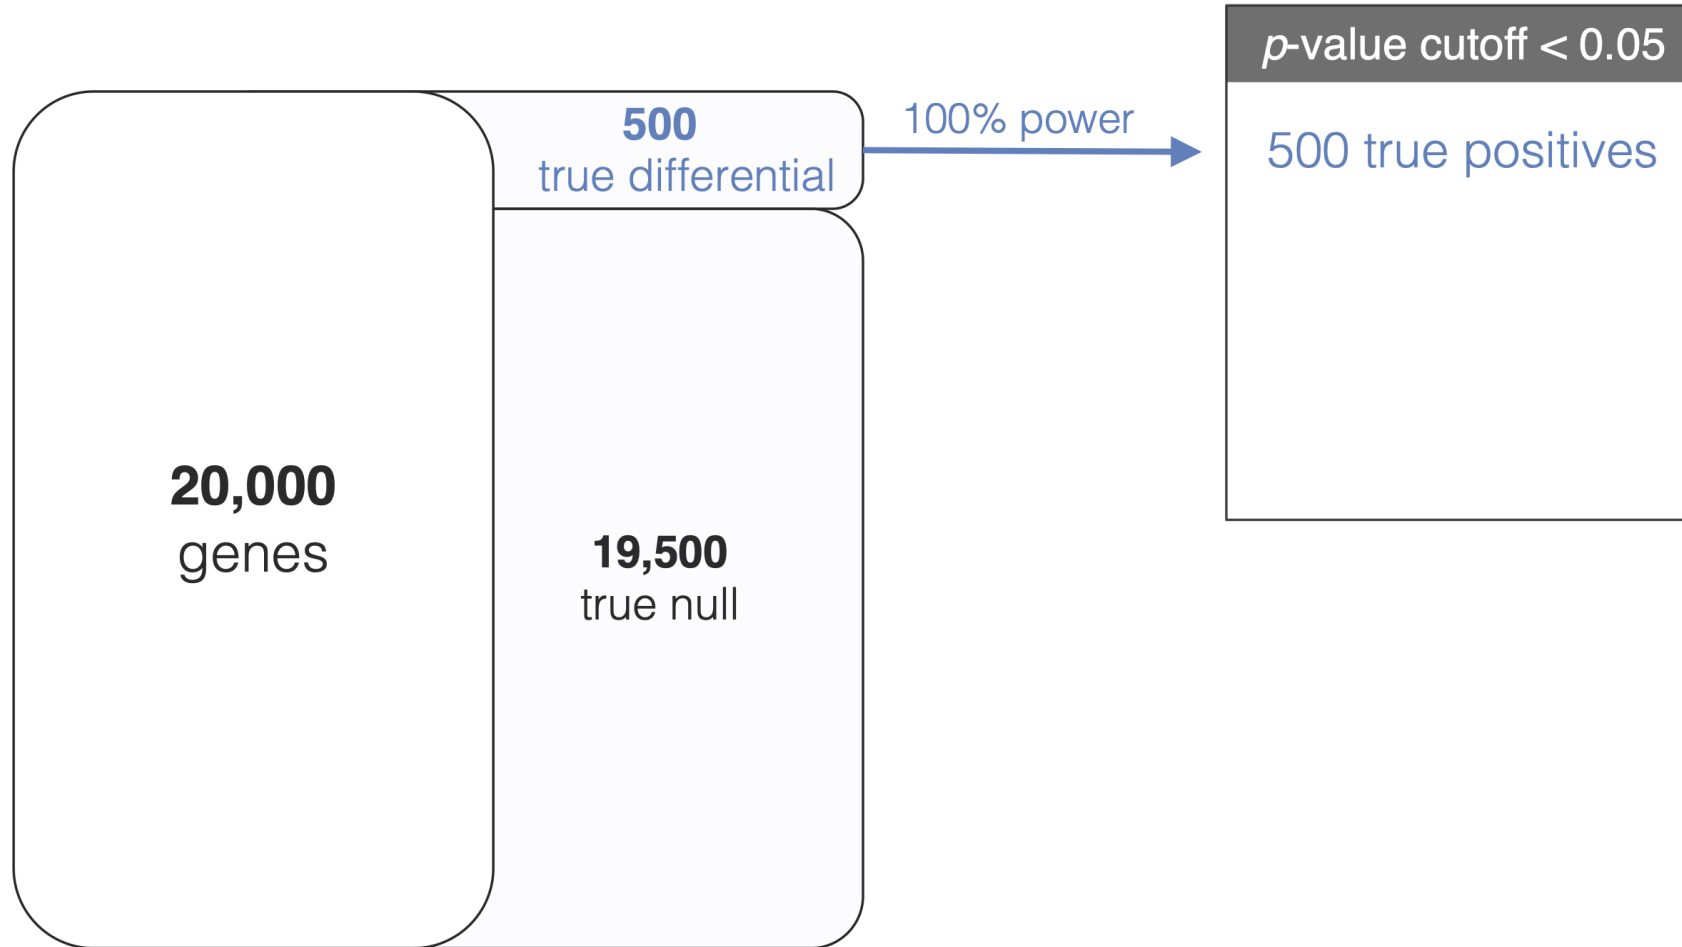
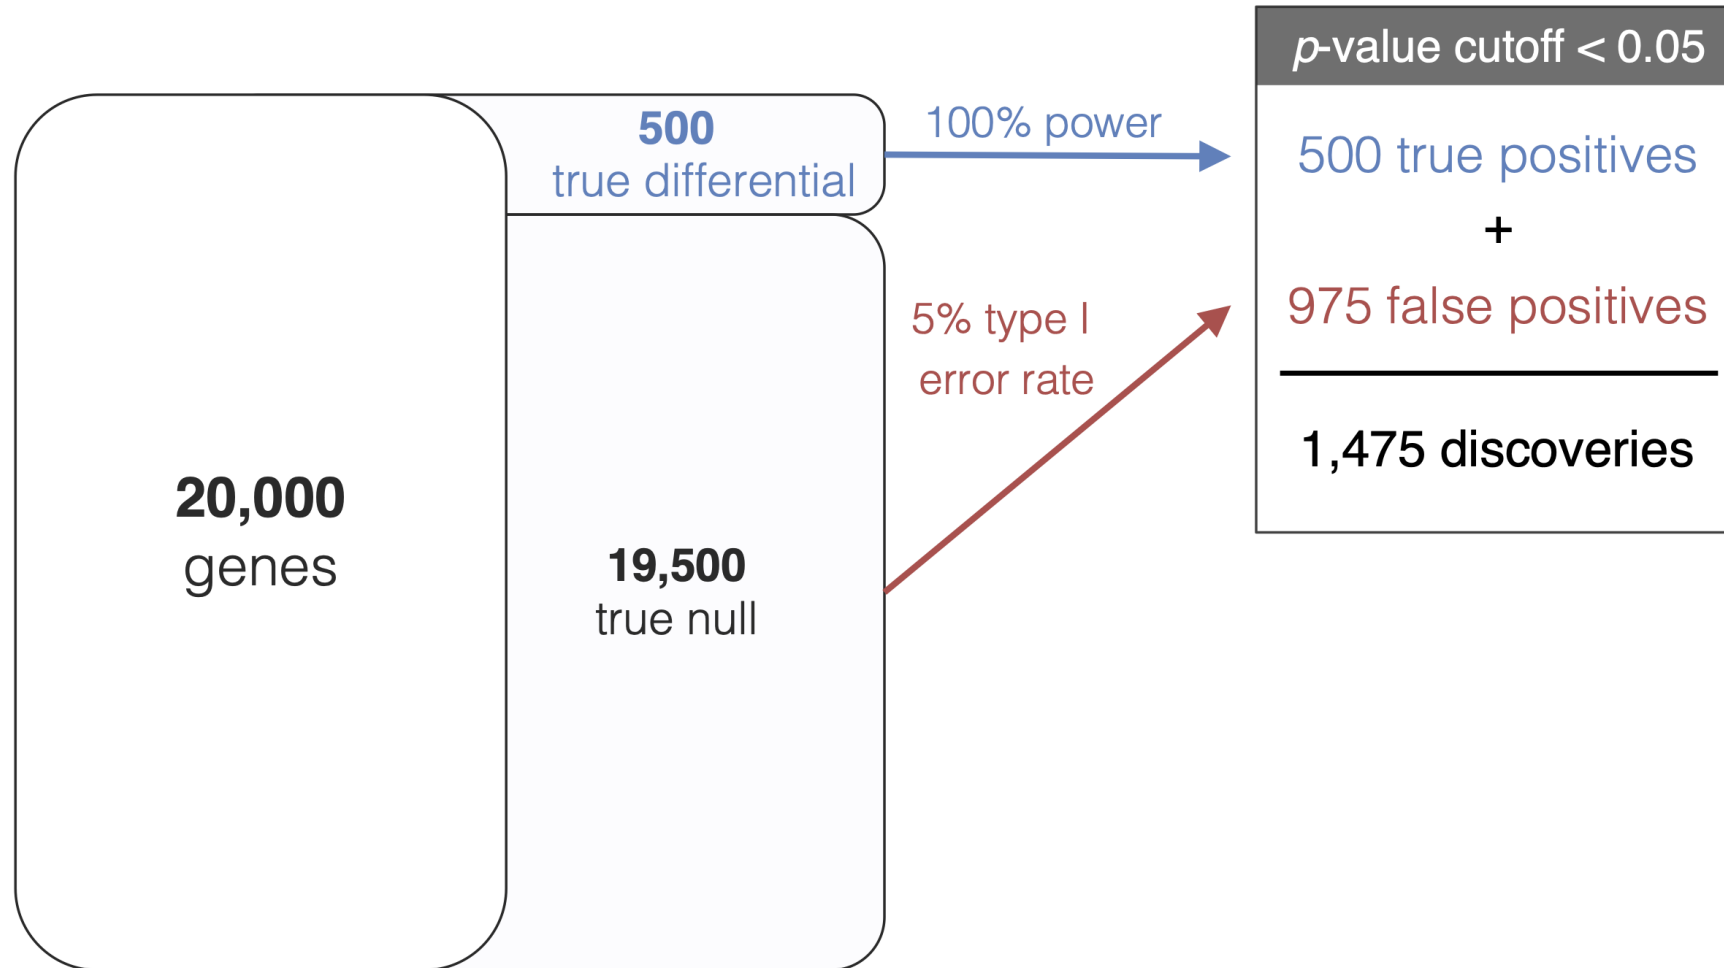# Multiple comparisons and error rates

20,000
genes

Multiple comparisons and error rates

Multiple comparisons and error rates

20,000
genes

500
true differential

19,500
true null

100% power

p-value cutoff < 0.05

500 true positives

Multiple comparisons and error rates

Multiple comparisons and error rates

20,000 genes

500 true differential — 100% power →

19,500 true null — 5% type I error rate ↗

p-value cutoff < 0.05

500 true positives

+

975 false positives

_____

1,475 discoveries

o best case scenario (100% power) yields 66% false positives

o problem worsens as power decreases and % null increases
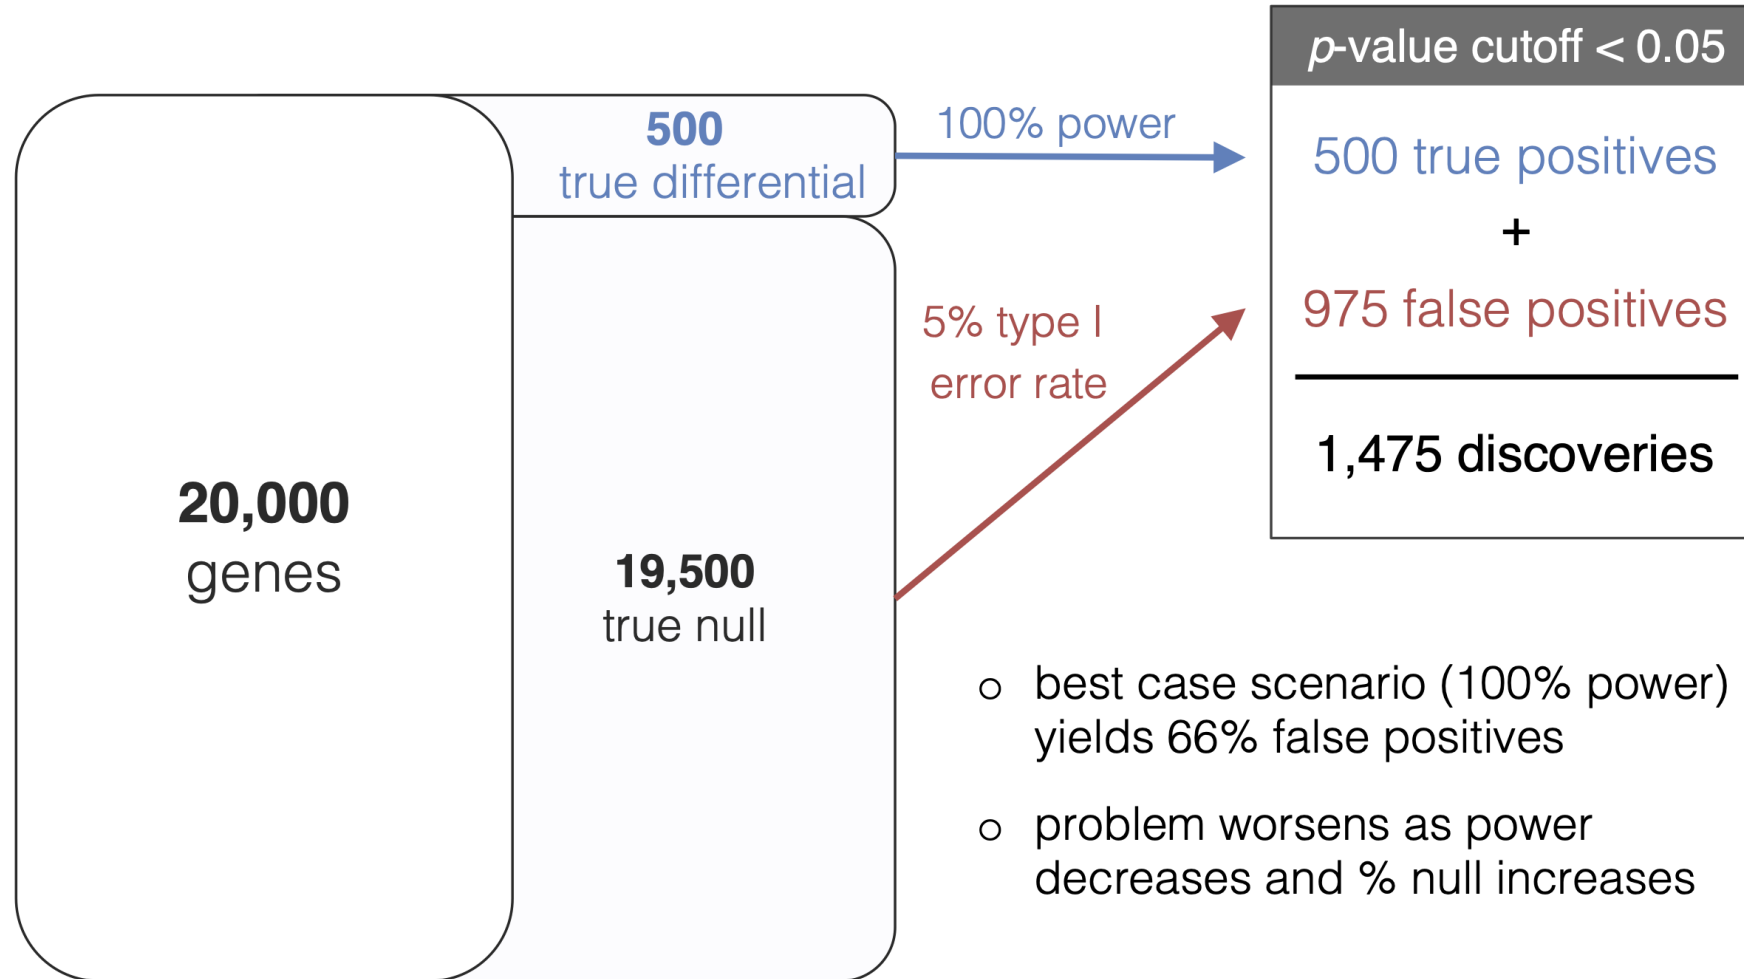
Family-Wise Error Rate (FWER)

- **FWER** is the probability of making at least one error when testing m tests

- **Control the FWER**: limit the probability of making at least one incorrect decision

- One example: the **Bonferroni** correction for $\alpha = 0.05$:

$$\text{If } P(\text{at least one error on } m \text{ tests}) < \alpha$$

$$\Rightarrow P(\text{at least one error on } m \text{ tests}) < \sum_{i=1}^{m} P(\text{error on test } i)$$

$$\sum_{i=1}^{m} P(\text{error on test } i) = m\alpha_{Bon}$$

$$\alpha_{Bon} = \frac{\alpha}{m} = \frac{0.05}{m}$$

Bonferroni correction: controlling the FWER

Can think of controlling the probability of at least one false positive in two ways:

1. **Adjust the p-values; keep same α:**

$$p_{Bon,i} = mp_i \text{ (more technically correct: } p_{Bon,i} = min(mp_i, 1))$$

Then, threshold $p_{Bon,i}$ at α

2. **Adjust the α threshold; keep same p-values:**

$$\alpha_{Bon} = \frac{\alpha}{m}$$

Then, threshold p_i at α_{Bon}

Multiple test correction is an active area of statistical research

- Bonferroni correction is very **conservative** (i.e. controls the FWER even lower than α in many settings)

- Several other options are better

- For example, the Holm procedure: multiplier for p-value correction is not the same for all genes; more powerful

$$p_{Holm,1} = mp_1$$

$$p_{Holm,2} = (m-1)p_2$$

$$p_{Holm,3} = (m-2)p_3$$

$$\vdots$$

$$\Rightarrow FWER \leq \alpha$$

How practical is the FWER in high-throughput biology?

- Why do we care so much about making one single error??

- One extreme way to ensure no Type I errors: reject no hypotheses! 😃

  - However, then our power is zero... 😭

- **Being overly strict about Type I error leads to greater Type II error** (loss of power)

Radical idea: it's OK to make multiple mistakes, as long as you also have some true positives!

Enter: the False Discovery Rate (FDR)

Controlling the False Discovery Rate: a Practical and Powerful Approach to Multiple Testing

By YOAV BENJAMINI† and YOSEF HOCHBERG

Tel Aviv University, Israel

SUMMARY

The common approach to the multiplicity problem calls for controlling the familywise error rate (FWER). This approach, though, has faults, and we point out a few. A different approach to problems of multiple significance testing is presented. It calls for controlling the expected proportion of falsely rejected hypotheses—the false discovery rate. This error rate is equivalent to the FWER when all hypotheses are true but is smaller otherwise. Therefore, in problems where the control of the false discovery rate rather than that of the FWER is desired, there is potential for a gain in power. A simple sequential Bonferroni-type procedure is proved to control the false discovery rate for independent test statistics, and a simulation study shows that the gain in power is substantial. The use of the new procedure and the appropriateness of the criterion are illustrated with examples.

Benjamini Y, Hochberg Y. "**Controlling the False Discovery Rate: a Practical and Powerful Approach to Multiple Testing.**" *Journal of the Royal Statistical Society: Series B (Methodological).* 1995 Jan;57(1): 289-300.

Over 70K citations!!

False Discovery Rate

	Null True	Alternative True	Total
Not Called Significant	U	T	$m - R$
Called Significant	V	S	R
	m_0	$m - m_0$	m

V = # Type I errors [false positives]

FDR is designed to control the expected proportion of false positives (V) among all hypotheses where the null has been rejected (R)

False Discovery Rate

	Null True	Alternative True	Total
Not Called Significant	U	T	m - R
Called Significant	V	S	R
	m_0	$m-m_0$	m

V = # Type I errors [false positives]

$$FDR = E\left[\frac{V}{R}\right]$$

FDR vs FPR vs FWER

- **False Discovery Rate (FDR)** is the rate that significant features (R) are truly null

$$FDR = E\left[\frac{V}{R}\right]$$

- **False Positive Rate (FPR)** is the rate that truly null features (m_0) are called significant

$$FPR = E\left[\frac{V}{m_0}\right]$$

- **Family-Wise Error Rate (FWER)** is the probability that the number of truly null features rejected (V) is at least 1

$$\text{FWER} = P(V \geq 1)$$

Benjamini-Hochberg FDR (BH procedure)

- Proposed the idea of controlling FDR instead of FWER

- Proposed a procedure for doing so

  - note that we know R, but we don't know V

- Procedure: control FDR at level q

  1. order the raw p-values $p_1 \leq p_2 \leq \ldots \leq p_m$

  2. find test with highest rank j such that $p_j < \frac{jq}{m}$

  3. declare all smaller ranks up to j significant

Controlling FDR at level $q = 0.05$

Rank (j)	P-value
1	0.0008
2	0.009
3	0.127
4	0.205
5	0.396
6	0.450
7	0.641
8	0.781
9	0.900
10	0.993

Controlling FDR at level $q = 0.05$

Rank (j)	P-value	(j/m)*q
1	0.0008	0.005
2	0.009	0.010
3	0.127	0.015
4	0.205	0.020
5	0.396	0.025
6	0.450	0.030
7	0.641	0.035
8	0.781	0.040
9	0.900	0.045
10	0.993	0.050

Controlling FDR at level $q = 0.05$

Rank (j)	P-value	(j/m)*q	Reject H_0?
1	0.0008	0.005	✓
2	0.009	0.010	✓
3	0.127	0.015	
4	0.205	0.020	
5	0.396	0.025	
6	0.450	0.030	
7	0.641	0.035	
8	0.781	0.040	
9	0.900	0.045	
10	0.993	0.050	

Controlling FDR at level $q = 0.05$

Rank (j)	P-value	(j/m)*q	Reject H_0?	$FWER_{Bon} < 0.05$?
1	0.0008	0.005	✓	✓
2	0.009	0.010	✓	
3	0.127	0.015		
4	0.205	0.020		
5	0.396	0.025		
6	0.450	0.030		
7	0.641	0.035		
8	0.781	0.040		
9	0.900	0.045		
10	0.993	0.050		

Where $\alpha_{Bon} = 0.05/10 = 0.005$

BH FDR values given in `limma` by default

```
topTable(ebFit, coef = "genotypeNrlKO")
```

```
##                    logFC  AveExpr          t      P.Value    adj.P.Val        B
## 1450946_at    -4.7134891 8.733103 -15.583683 5.390931e-18 2.431364e-13 28.40938
## 1442070_at    -0.9669400 8.268122  -9.787407 6.896220e-12 1.555132e-07 16.46086
## 1443457_at    -1.1530057 6.049428  -9.091412 4.947029e-11 6.029428e-07 14.68843
## 1422679_s_at  -2.1187394 9.495435  -9.002625 6.388954e-11 6.029428e-07 14.45703
## 1431708_a_at  -2.1610625 8.591450  -8.967950 7.062007e-11 6.029428e-07 14.36634
## 1433050_at    -1.5083887 8.468891  -8.923932 8.021234e-11 6.029428e-07 14.25096
## 1426288_at    -4.2100657 9.323796  -8.772209 1.246491e-10 8.031144e-07 13.85106
## 1418108_at     0.8662091 8.260647   8.537984 2.475294e-10 1.395478e-06 13.22713
## 1450770_at     1.3103576 7.357033   8.128847 8.332369e-10 3.747976e-06 12.11870
## 1457802_at    -0.8362030 7.778189  -8.117112 8.629984e-10 3.747976e-06 12.08657
```

Or, obtain them yourself for any vector of p-values p with `p.adjust(p, method="BH")`

Other ways to control FDR

- BH is just one (the first) method to control FDR

- Since the publication of the BH method, other methods have been proposed

- One of the most popular is Storey's q-value

Statistical significance for genomewide studies

John D. Storey*[†] and Robert Tibshirani[‡]

*Department of Biostatistics, University of Washington, Seattle, WA 98195; and ‡Departments of Health Research and Policy and Statistics, Stanford University, Stanford, CA 94305

Edited by Philip P. Green, University of Washington School of Medicine, Seattle, WA, and approved May 30, 2003 (received for review January 28, 2003)

With the increase in genomewide experiments and the sequencing of multiple genomes, the analysis of large data sets has become commonplace in biology. It is often the case that thousands of features in a genomewide data set are tested against some null hypothesis, where a number of features are expected to be significant. Here we propose an approach to measuring statistical significance in these genomewide studies based on the concept of the false discovery rate. This approach offers a sensible balance between the number of true and false positives that is automatically calibrated and easily interpreted. In doing so, a measure of statistical significance called the *q* value is associated with each tested feature. The *q* value is similar to the well known *p* value, except it is a measure of significance in terms of the false discovery rate rather than the false positive rate. Our approach avoids a flood of false positive results, while offering a more liberal criterion than what has been used in genome scans for linkage.

false discovery rates | genomics | multiple hypothesis testing | *q* values

to the method in ref. 5 under certain assumptions. Also, ideas similar to FDRs have appeared in the genetics literature (1, 13).

Similarly to the *p* value, the *q* value gives each feature its own individual measure of significance. Whereas the *p* value is a measure of significance in terms of the false positive rate, the *q* value is a measure in terms of the FDR. The false positive rate and FDR are often mistakenly equated, but their difference is actually very important. Given a rule for calling features significant, the false positive rate is the rate that truly null features are called significant. The FDR is the rate that significant features are truly null. For example, a false positive rate of 5% means that on average 5% of the truly null features in the study will be called significant. A FDR of 5% means that among all features called significant, 5% of these are truly null on average.

The *q* value provides a measure of each feature's significance, automatically taking into account the fact that thousands are simultaneously being tested. Suppose that features with *q* values ≤5% are called significant in some genomewide test of significance. This results in a FDR of 5% among the significant features. A

- `qvalue` package implementation: provides adjusted p-values

Storey's q-value vs BH (Conceptual)

- Just like BH, is focused on the proportion of discoveries that are false positives

- *Conceptual* difference between BH and Storey's q-value is:

  - BH **controls** the FDR

  - q-values give an unbiased **estimate** of the FDR (will control the FDR on eaverage)

Storey's q-value vs BH (Mathematical)

- Mathematically, the difference between the two is in how m_0 is estimated

  - Or equivalently, how $\pi_0 = \frac{m_0}{m}$ is estimated (since m is known)

  - π_0 represents the proportion of tests that are truly null

- q-value:

$$\hat{q}(p_i) = \min_i \left(\frac{\hat{\pi}_0 m}{rank(p_i)} p_i, \; 1 \right)$$

- q-value and BH-adjusted p-values are equivalent when $\pi_0 = 1$

$$\hat{p}_{BH}(p_i) = \min_i \left(\frac{m}{rank(p_i)} p_i, \; 1 \right)$$

(BH conservatively assumes that $\pi_0 = 1$)

BH vs q-value in our example

Rank (j)	P-value	$\hat{p}_{BH}(p_i)$	$\hat{q}(p_i)$
1	0.0008	0.008	0.008
2	0.009	0.045	0.045
3	0.127	0.423	0.423
4	0.205	0.513	0.513
5	0.396	0.792	0.750
6	0.450	0.750	0.750
7	0.641	0.916	0.916
8	0.781	0.976	0.976
9	0.900	1.000	0.993
10	0.993	0.993	0.993

# Compounding issues of multiple comparisons

- What if you're not only testing 45K genes, but also multiple tests per gene (e.g. multiple contrasts, such as several two-group comparisons)?

- Classical procedures for adjustment:

    - Tukey multiple comparison procedure
    - Scheffe multiple comparison procedure
    - Bonferroni or Holm FWER correction

- In our setting, we can also apply BH to all p-values globally

    - `limma::decideTests(pvals, method="global")` for a matrix of p-values or `eBayes` output (e.g. rows = genes, columns = contrasts)

    - p-values are combined, adjusted globally, then separated back out and sorted

Assumptions about p-values

- Implicit assumption for all multiple testing correction methods:

<div align="center">

p-value distribution is "well-behaved"

</div>

- What does this mean?

  - primarily, that the distribution of p-values under the null is **uniform**

Histogram of pval

p-value distributions

Spike of small p-values indicates non-null tests:

Histogram of pval

Great primer on how things can go wrong: http://varianceexplained.org/statistics/interpreting-pvalue-histogram/

What if p-values are *poorly behaved*?

- FDR estimates can be invalid (assumptions are violated)

- Solution: compute p-values "empirically" using resampling/permutation/bootstrap techniques

- **Bootstrap**: take repeated random samples *with replacement* from your dataset (randomly assigning covariate of interest) and compute statistic; repeat many times and use bootstrap statistics as your sampling distribution rather than a t, Normal, F, χ^2, etc

- **Permutation**: construct a simulated version of your dataset that satisfies the null hypothesis and compute statistic (e.g. shuffle group labels for a two-group comparison); repeat many times and use permutation statistics as your sampling distribution rather than a t, Normal, F, χ^2, etc

- Downsides: often computationally intensive for genomics; not optimal for small sample sizes