

# Statistical Methods for High Dimensional Biology

More limma & multiple testing

Keegan Korthauer

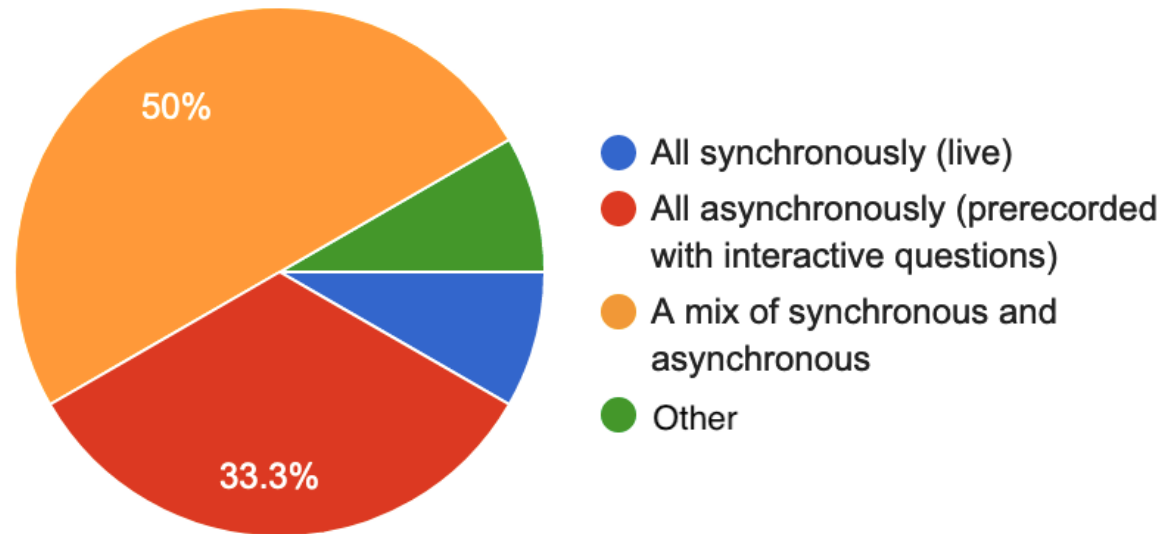
8 February 2021

with slide contributions from Gabriela Cohen Freue, Jenny Bryan, and Sara Mostafavi

# Announcements

- Feedback posted on your initial project proposals (see the issue in your team's repository)
- Reminder that the **finalized project proposals** are due Feb 18
- Instructions on how to submit your deliverables have been updated for clarity on the **website**
  - in particular pay attention to the recommended way to grab the link to the submitted files from your latest commit

# Course format feedback



- Preference for mix of synchronous & asynchronous
- 😊 Positive feedback on interactive quiz questions / polls / chat
- 😞 Several people felt a lack of communication - let's utilize the **Discussion repo** more!

# Learning objectives

- How to use `limma` to perform genome-wide differential expression testing on microarray data
- Understand the key differences between `limma` and `lm`
- Explain why multiple testing increases the number of errors we make by chance
- Be able to adjust for multiple comparisons by controlling the False Discovery Rate
  - e.g. using Benjamini-Hochberg or Storey's q-value

# Recall: The hybrid estimator in `limma`

$$\tilde{s}_g^2 = \frac{d_0 s_0^2 + d s_g^2}{d_0 + d}$$

- recall that  $(s_0, d_0)$  are the *prior* parameters for  $\sigma_g^2$  (random variable):

$$\frac{1}{\sigma_g^2} \sim \frac{1}{d_0 s_0^2} \chi_{d_0}^2$$

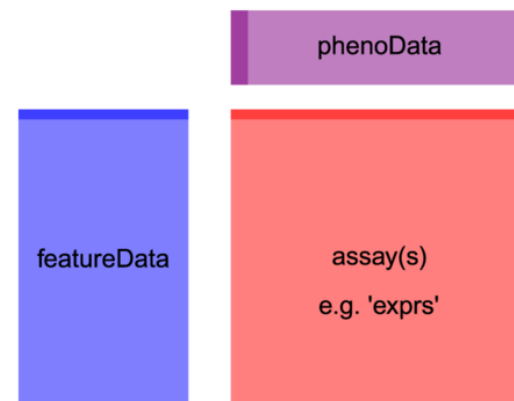
- the prior parameters incorporate information from all genes which allows us to **shrink/nudge the gene-specific variances toward a common consensus**
  - they are estimated from the data - the formulas for  $(s_0, d_0)$  and their derivations are beyond the scope of the course (but `limma` takes care of the details for us)
  - note that  $(s_0, d_0)$  do not depend on  $g$

# limma quickstart

$$\mathbf{Y}_g = \mathbf{X}\boldsymbol{\alpha}_g + \boldsymbol{\epsilon}_g,$$

$$E(\boldsymbol{\epsilon}_g) = 0, \text{Var}(\boldsymbol{\epsilon}_g) = \sigma_g^2, \epsilon_{ig} \perp \epsilon_{jg},$$

- Within each gene observations are iid / constant variance
- `lmFit()` function in `limma` carries out multiple linear regression on each gene
- Usage: `lmFit(myDat, desMat)`
  - `myDat` is a data frame, matrix with features in rows and samples in columns ( $G$  genes by  $N$  samples), or `ExpressionSet` object
  - `desMat` is a design matrix (output of `model.matrix(y ~ x)`;  $N$  samples by  $p$  parameters)



Bioconductor's ExpressionSet class


Let's run `l i m m a` for the interactive model with age

$$y_{ig} = \theta + \tau_{KO}x_{ig,KO} + \tau_{Age}x_{ig,Age} + \tau_{KO:Age}x_{ig,KO}x_{ig,Age}$$

- $i$  indexes mouse,  $g$  indexes genes
- $x_{ig,KO}$  is the dummy/indicator variable for the NrlKO group
- $x_{ig,Age}$  is the continuous age variable

# Let's run `limma` for the interactive model with age

Example gene (but we want to fit this model on all genes):





# Arranging input for `limFit`: Bioconductor way

```
eSet
```

```
## ExpressionSet (storageMode: lockedEnvironment)
## assayData: 45101 features, 39 samples
##   element names: exprs
## protocolData: none
## phenoData
##   sampleNames: GSM92610 GSM92611 ... GSM92648 (39 total)
##   varLabels: sample_id dev_stage genotype age
##   varMetadata: labelDescription
## featureData: none
## experimentData: use 'experimentData(object)'
##   pubMedIds: 16505381
## Annotation: GPL1261
```

# Arranging input for limFit: Separate expression & metadata

```
myDat <- assayData(eset)$exprs  
myDat[1:5,1:5]
```

```
##           GSM92610 GSM92611  GSM92612 GSM92613  GSM92614  
## 1415670_at    7.108863 7.322392   7.420947 7.351444   7.240428  
## 1415671_at    9.714002 9.797742   9.831072 9.658442   9.708906  
## 1415672_at    9.429030 9.846977  10.003092 9.914112  10.174548  
## 1415673_at    8.426974 8.404206   8.594600 8.404206   8.835334  
## 1415674_a_at  8.498338 8.458287   8.426651 8.372776   8.541722
```

```
myMeta <- pData(eset)  
myMeta[1:5,]
```

```
##           sample_id dev_stage genotype age  
## GSM92610    GSM92610      4W      Nr1k0  28  
## GSM92611    GSM92611      4W      Nr1k0  28  
## GSM92612    GSM92612      4W      Nr1k0  28  
## GSM92613    GSM92613      4W      Nr1k0  28  
## GSM92614    GSM92614     E16      Nr1k0  -4
```

# Formulating input for `lme4`: Design Matrix

Bioconductor way:

```
desMat <- model.matrix(~ genotype*age,  
                        data = pData(eset))
```

Equivalently, if using the separate way:

```
desMat <- model.matrix(~ genotype*age,  
                        data = myMeta)
```

desMat

	(Intercept)	genotypeNr1k0	age	genotypeNr1k0:age
## GSM92610	1	1	28	28
## GSM92611	1	1	28	28
## GSM92612	1	1	28	28
## GSM92613	1	1	28	28
## GSM92614	1	1	-4	-4
## GSM92615	1	1	-4	-4
## GSM92616	1	1	-4	-4
## GSM92617	1	1	10	10
## GSM92618	1	1	10	10
## GSM92619	1	1	10	10
## GSM92620	1	1	10	10
## GSM92621	1	1	2	2
## GSM92622	1	1	2	2
## GSM92623	1	1	2	2
## GSM92624	1	1	2	2
## GSM92625	1	1	6	6
## GSM92626	1	1	6	6
## GSM92627	1	1	6	6
## GSM92628	1	1	6	6
## GSM92629	1	0	28	0
## GSM92630	1	0	28	0
## GSM92631	1	0	28	0
## GSM92632	1	0	28	0
## GSM92633	1	0	-4	0
## GSM92634	1	0	-4	0
## GSM92635	1	0	-4	0
## GSM92636	1	0	-4	0
## GSM92637	1	0	10	0
## GSM92638	1	0	10	0

# Computation is fast

```
system.time(gFit <- lmFit(eset,  
                          model.matrix(~ genotype*age, data = pData(eset))))
```

```
##      user  system elapsed  
##    0.111    0.025    0.136
```

- Equivalently, using the 'separate' way:

```
gFit <- lmFit(myDat,  
              model.matrix(~ genotype*age, data = myMeta))
```

- Using `lmFit` to fit an interactive model on 45K probesets takes a fraction of a second
- The time-intensive parts of an analysis lie in selecting the model and covariates, choosing how to parameterize it, and interpreting the output

# Output of `lmFit`

```
summary(gFit)
```

```
##              Length Class  Mode
## coefficients 180404 -none- numeric
## rank         1 -none- numeric
## assign       4 -none- numeric
## qr           5 qr      list
## df.residual  45101 -none- numeric
## sigma        45101 -none- numeric
## cov.coefficients 16 -none- numeric
## stdev.unscaled 180404 -none- numeric
## pivot         4 -none- numeric
## Amean         45101 -none- numeric
## method        1 -none- character
## design       156 -none- numeric
```

```
nrow(eset)
```

```
## Features
##      45101
```

```
nrow(eset)*4
```

```
## Features
##      180404
```

# Output of `lmFit`

```
summary(gFit)
```

```
##              Length Class  Mode
## coefficients 180404 -none- numeric
## rank         1 -none- numeric
## assign       4 -none- numeric
## qr           5 qr      list
## df.residual  45101 -none- numeric
## sigma       45101 -none- numeric
## cov.coefficients 16 -none- numeric
## stdev.unscaled 180404 -none- numeric
## pivot        4 -none- numeric
## Amean       45101 -none- numeric
## method       1 -none- character
## design      156 -none- numeric
```

```
nrow(eset)
```

```
## Features
##      45101
```

```
nrow(eset)*4
```

```
## Features
##      180404
```

- OK... but where are the shrunken variable estimates?? How do I pull out p-values??
- Actually, we haven't carried out the empirical Bayesian computation yet -- need to run `eBayes()`!

# eBayes()

```
summary(gFit)
```

##	Length	Class	Mode
## coefficients	180404	-none-	numeric
## rank	1	-none-	numeric
## assign	4	-none-	numeric
## qr	5	qr	list
## df.residual	45101	-none-	numeric
## sigma	45101	-none-	numeric
## cov.coefficients	16	-none-	numeric
## stdev.unscaled	180404	-none-	numeric
## pivot	4	-none-	numeric
## Amean	45101	-none-	numeric
## method	1	-none-	character
## design	156	-none-	numeric

```
summary(ebFit <- eBayes(gFit))
```

##	Length	Class	Mode
## coefficients	180404	-none-	numeric
## rank	1	-none-	numeric
## assign	4	-none-	numeric
## qr	5	qr	list
## df.residual	45101	-none-	numeric
## sigma	45101	-none-	numeric
## cov.coefficients	16	-none-	numeric
## stdev.unscaled	180404	-none-	numeric
## pivot	4	-none-	numeric
## Amean	45101	-none-	numeric
## method	1	-none-	character
## design	156	-none-	numeric
## df.prior	1	-none-	numeric
## s2.prior	1	-none-	numeric
## var.prior	4	-none-	numeric
## proportion	1	-none-	numeric
## s2.post	45101	-none-	numeric
## t	180404	-none-	numeric
## df.total	45101	-none-	numeric
## p.value	180404	-none-	numeric
## lods	180404	-none-	numeric
## F	45101	-none-	numeric
## F.p.value	45101	-none-	numeric

# Components of the empirical Bayes estimators

math	plain english	limma	numerical result	also in lm?
$s_g^2$	gene-specific residual variance	<code>gFit\$sigma^2</code>	45K numbers	✓
$d$	residual degrees of freedom ( $n - p$ )	<code>gFit\$df.residual</code>	$39 - 4 = 35^*$	✓
$s_0^2$	mean of inverse $\chi^2$ prior for $s_g^2$	<code>ebFit\$s2.prior</code>	0.048	
$d_0$	degrees of freedom for the prior	<code>ebFit\$df.prior</code>	2.61	
$\tilde{s}_g^2$	posterior mean of $s_g^2$ (i.e. moderated residual variance)	<code>ebFit\$s2.post</code>	45K numbers	

\* limma can handle more complicated models where this is not the same for each gene, so this is actually a vector of 45K copies of the number 35



# topTable() will help us extract relevant output in a convenient format!

```
topTable(fit, coef=NULL, number=10, genelist=fit$genes, adjust.method="BH",  
        sort.by="B", resort.by=NULL, p.value=1, lfc=0, confint=FALSE)
```

fit	list containing a linear model fit produced by <code>lmFit</code> , <code>lm.series</code> , <code>gls.series</code> or <code>mr1m</code> . For <code>topTable</code> , <code>fit</code> should be an object of class <code>MArrayLM</code> as produced by <code>lmFit</code> and <code>eBayes</code> .
coef	column number or column name specifying which coefficient or contrast of the linear model is of interest. For <code>topTable</code> , can also be a vector of column subscripts, in which case the gene ranking is by F-statistic for that set of contrasts.
number	maximum number of genes to list
adjust.method	method used to adjust the p-values for multiple testing. Options, in increasing conservatism, include "none", "BH", "BY" and "holm". See <code>p.adjust</code> for the complete list of options. A NULL value will result in the default adjustment method, which is "BH".
sort.by	character string specifying statistic to rank genes by. Possible values for <code>topTable</code> and <code>topTableF</code> are "logFC", "AveExpr", "t", "P", "p", "B" or "none". (Permitted synonyms are "M" for "logFC", "A" or "Amean" for "AveExpr", "T" for "t" and "p" for "P".) Possibilities for <code>topTableF</code> are "F" or "none". Possibilities for <code>topTreat</code> are as for <code>topTable</code> except for "B".

... (truncated - see `?topTable` for full listing)

# Summary of `topTable` function

- `coef` is the argument where you specify the coefficient(s) you want to test for equality with zero (default is NULL; must be specified)
- `number` lets you control size of hit list (default is 10)
- `p.value` lets you specify a minimum adjusted p-value cutoff (default is 1)
- `lfc` lets you specify a minimum observed effect size - log2 fold change (default is 0)
- `sort.by` and `resort.by` give control over the ordering (default is by "B": log-odds that the gene is differentially expressed)
- `adjust.method` specifies how/if to adjust p-values for multiple testing (default is BH)


# topTable in action: genotypeNr1K0

```
topTable(ebFit, coef = "genotypeNr1K0")
```

##		logFC	AveExpr	t	P.Value	adj.P.Val	B
##	1450946_at	-4.7134891	8.733103	-15.583683	5.390931e-18	2.431364e-13	28.40938
##	1442070_at	-0.9669400	8.268122	-9.787407	6.896220e-12	1.555132e-07	16.46086
##	1443457_at	-1.1530057	6.049428	-9.091412	4.947029e-11	6.029428e-07	14.68843
##	1422679_s_at	-2.1187394	9.495435	-9.002625	6.388954e-11	6.029428e-07	14.45703
##	1431708_a_at	-2.1610625	8.591450	-8.967950	7.062007e-11	6.029428e-07	14.36634
##	1433050_at	-1.5083887	8.468891	-8.923932	8.021234e-11	6.029428e-07	14.25096
##	1426288_at	-4.2100657	9.323796	-8.772209	1.246491e-10	8.031144e-07	13.85106
##	1418108_at	0.8662091	8.260647	8.537984	2.475294e-10	1.395478e-06	13.22713
##	1450770_at	1.3103576	7.357033	8.128847	8.332369e-10	3.747976e-06	12.11870
##	1457802_at	-0.8362030	7.778189	-8.117112	8.629984e-10	3.747976e-06	12.08657

- `topTable(ebFit, coef = 2)` is equivalent here, but much less informative!!
- What is the null hypothesis here?
  - $H_0 : \tau_{KO} = 0$
  - this finds genes where the knockouts differ from the wild types *when age is zero*

# Plotting the top 6 probes for `genotypeNr1KO`



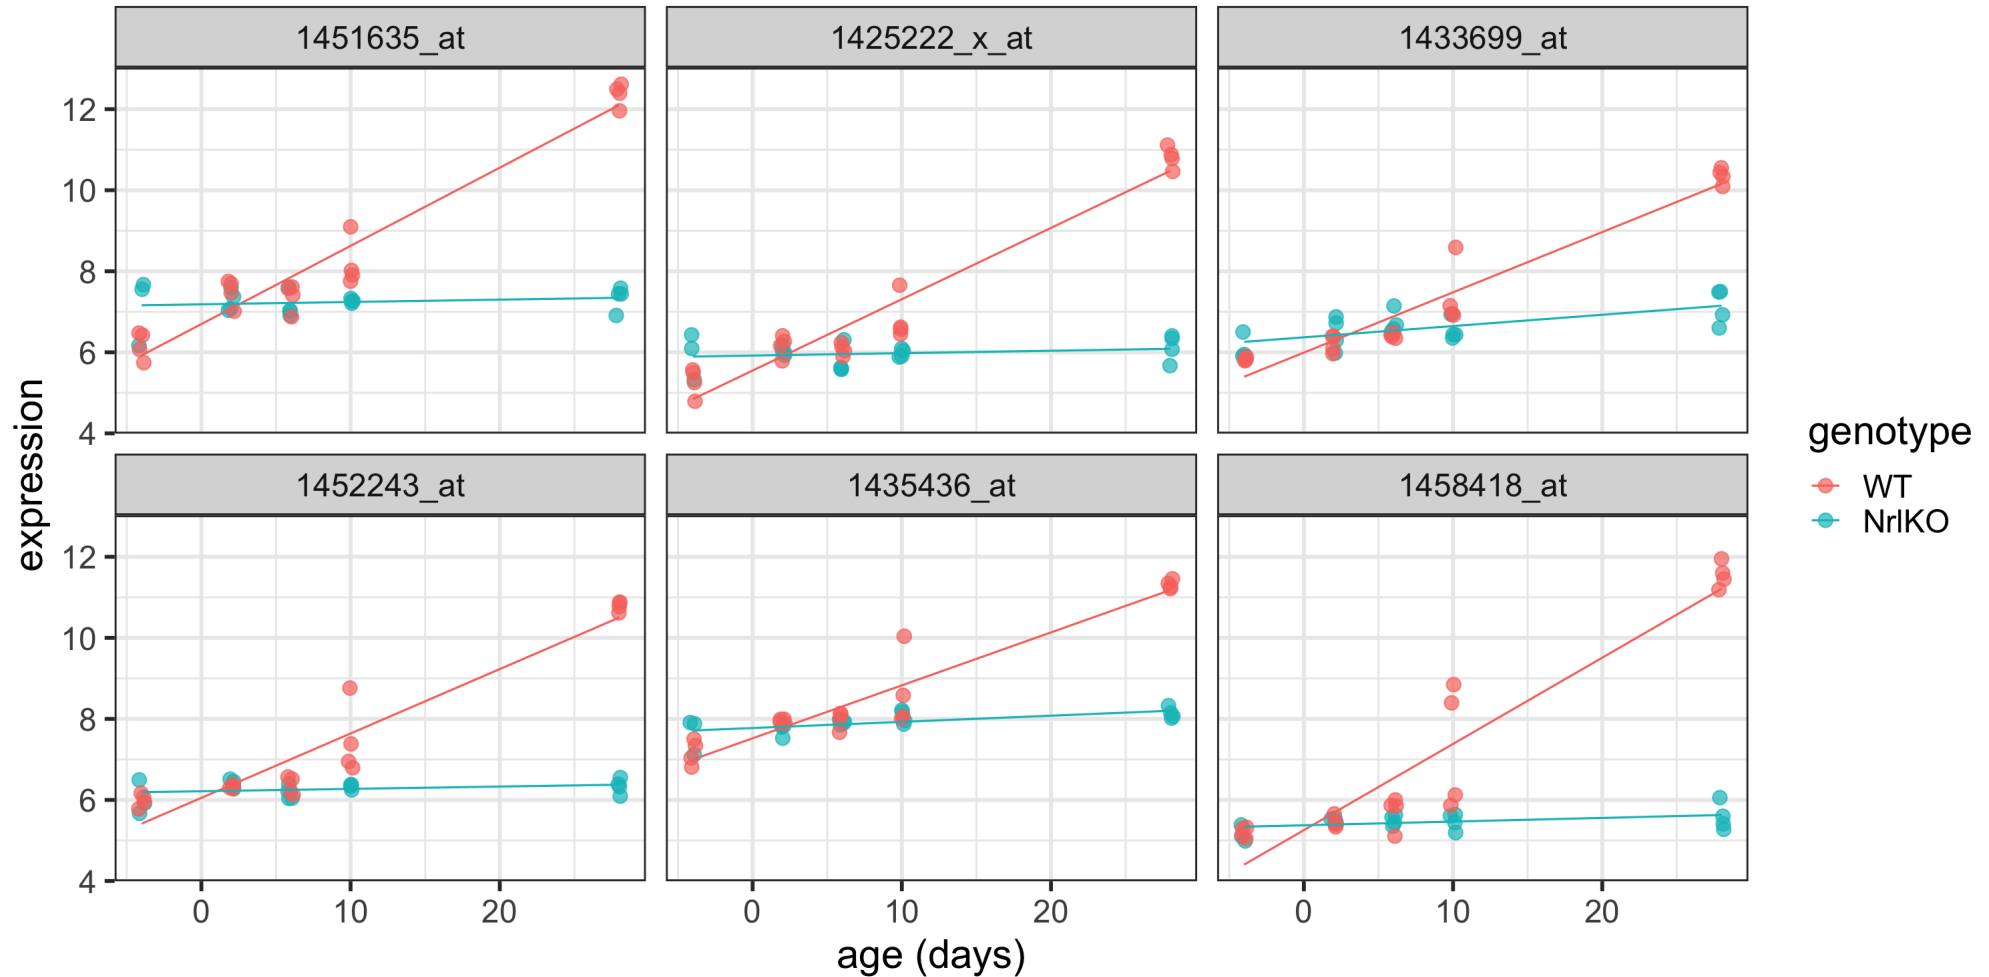
# topTable in action: age

```
topTable(ebFit, coef = "age")
```

##		logFC	AveExpr	t	P.Value	adj.P.Val	B
##	1451635_at	0.19294061	7.791237	20.78935	3.204717e-22	1.445359e-17	40.60768
##	1425222_x_at	0.17604111	6.514475	19.72633	1.962954e-21	4.426561e-17	38.79502
##	1433699_at	0.14875031	6.939973	17.80467	6.442530e-20	7.957605e-16	35.29412
##	1452243_at	0.15889179	6.840975	17.75658	7.057586e-20	7.957605e-16	35.20253
##	1435436_at	0.13067119	8.274676	17.61808	9.187120e-20	8.286966e-16	34.93763
##	1458418_at	0.21277027	6.270187	16.43686	9.326171e-19	7.006477e-15	32.60718
##	1424977_at	0.07814988	6.505464	16.36079	1.087456e-18	7.006477e-15	32.45260
##	1431174_at	0.16158579	7.418000	16.01901	2.183055e-18	1.230725e-14	31.75106
##	1421818_at	0.11250919	7.982278	15.66666	4.531133e-18	2.270651e-14	31.01562
##	1419069_at	0.09325234	8.222096	15.54775	5.813459e-18	2.621928e-14	30.76457

- `topTable(ebFit, coef = 3)` is equivalent here, but much less informative!!
- What is the null hypothesis here?
  - $H_0 : \tau_{Age} = 0$
  - this finds genes where age significantly affects gene expression *for WT*

# Plotting the top 6 genes for age



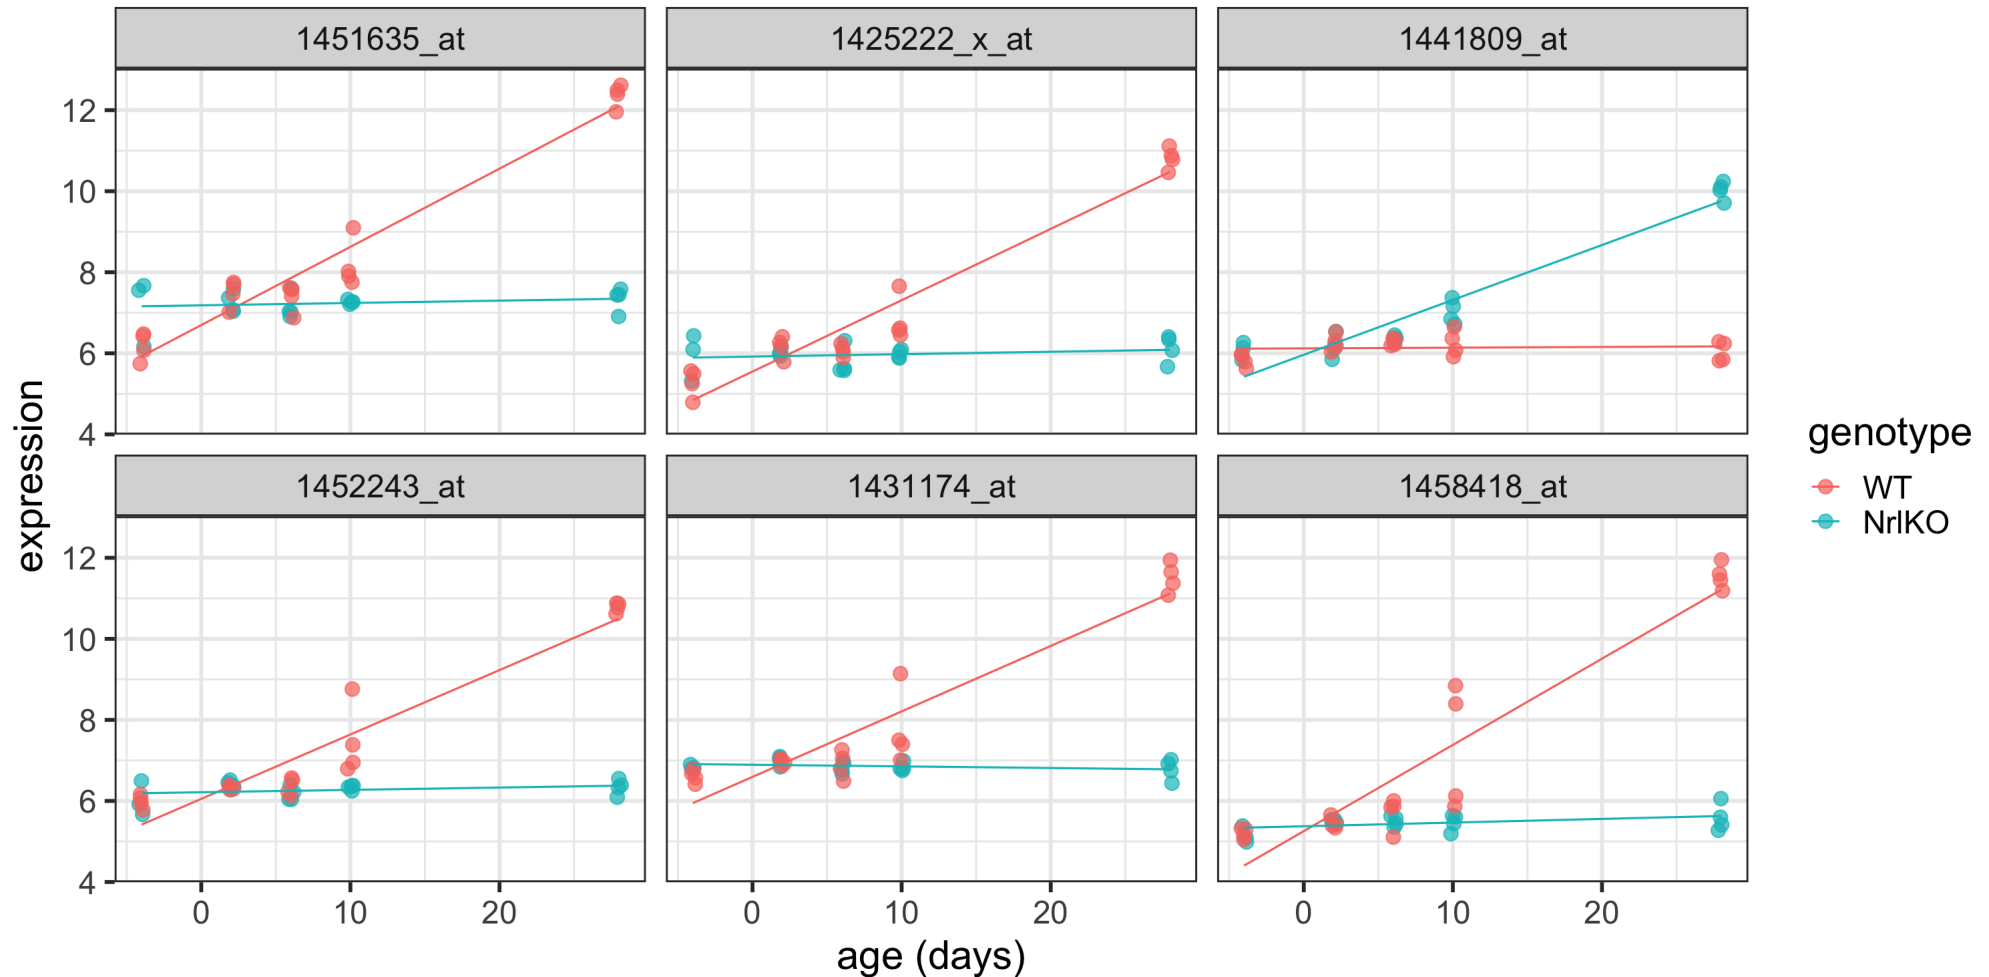
# topTable in action: genotypeNr1K0:age

```
topTable(ebFit, coef = "genotypeNr1K0:age")
```

##		logFC	AveExpr	t	P.Value	adj.P.Val	B
##	1451635_at	-0.1871800	7.791237	-14.00463	1.684302e-16	7.596369e-12	27.52010
##	1425222_x_at	-0.1700908	6.514475	-13.23451	9.961393e-16	2.246344e-11	25.74877
##	1441809_at	0.1336356	6.649440	12.55702	5.037221e-15	7.572790e-11	24.13110
##	1452243_at	-0.1531326	6.840975	-11.88284	2.670693e-14	3.011273e-10	22.46414
##	1431174_at	-0.1655894	7.418000	-11.39882	9.161231e-14	8.263614e-10	21.23123
##	1458418_at	-0.2037314	6.270187	-10.92850	3.122990e-13	2.347500e-09	20.00384
##	1435436_at	-0.1154500	8.274676	-10.80855	4.289363e-13	2.763637e-09	19.68613
##	1446715_at	-0.1409774	5.613122	-10.57909	7.912260e-13	4.325169e-09	19.07307
##	1442190_at	-0.1323038	5.896077	-10.54672	8.630965e-13	4.325169e-09	18.98601
##	1416306_at	0.1702059	6.560050	10.41285	1.238250e-12	5.584630e-09	18.62456

- `topTable(ebFit, coef = 4)` is equivalent here, but much less informative!!
- What is the null hypothesis here?
  - $H_0 : \tau_{KO:Age} = 0$
  - this finds genes where the effect of age is significantly different in each genotype

# Plotting the top 6 genes for `genotypeNr1K0:age`





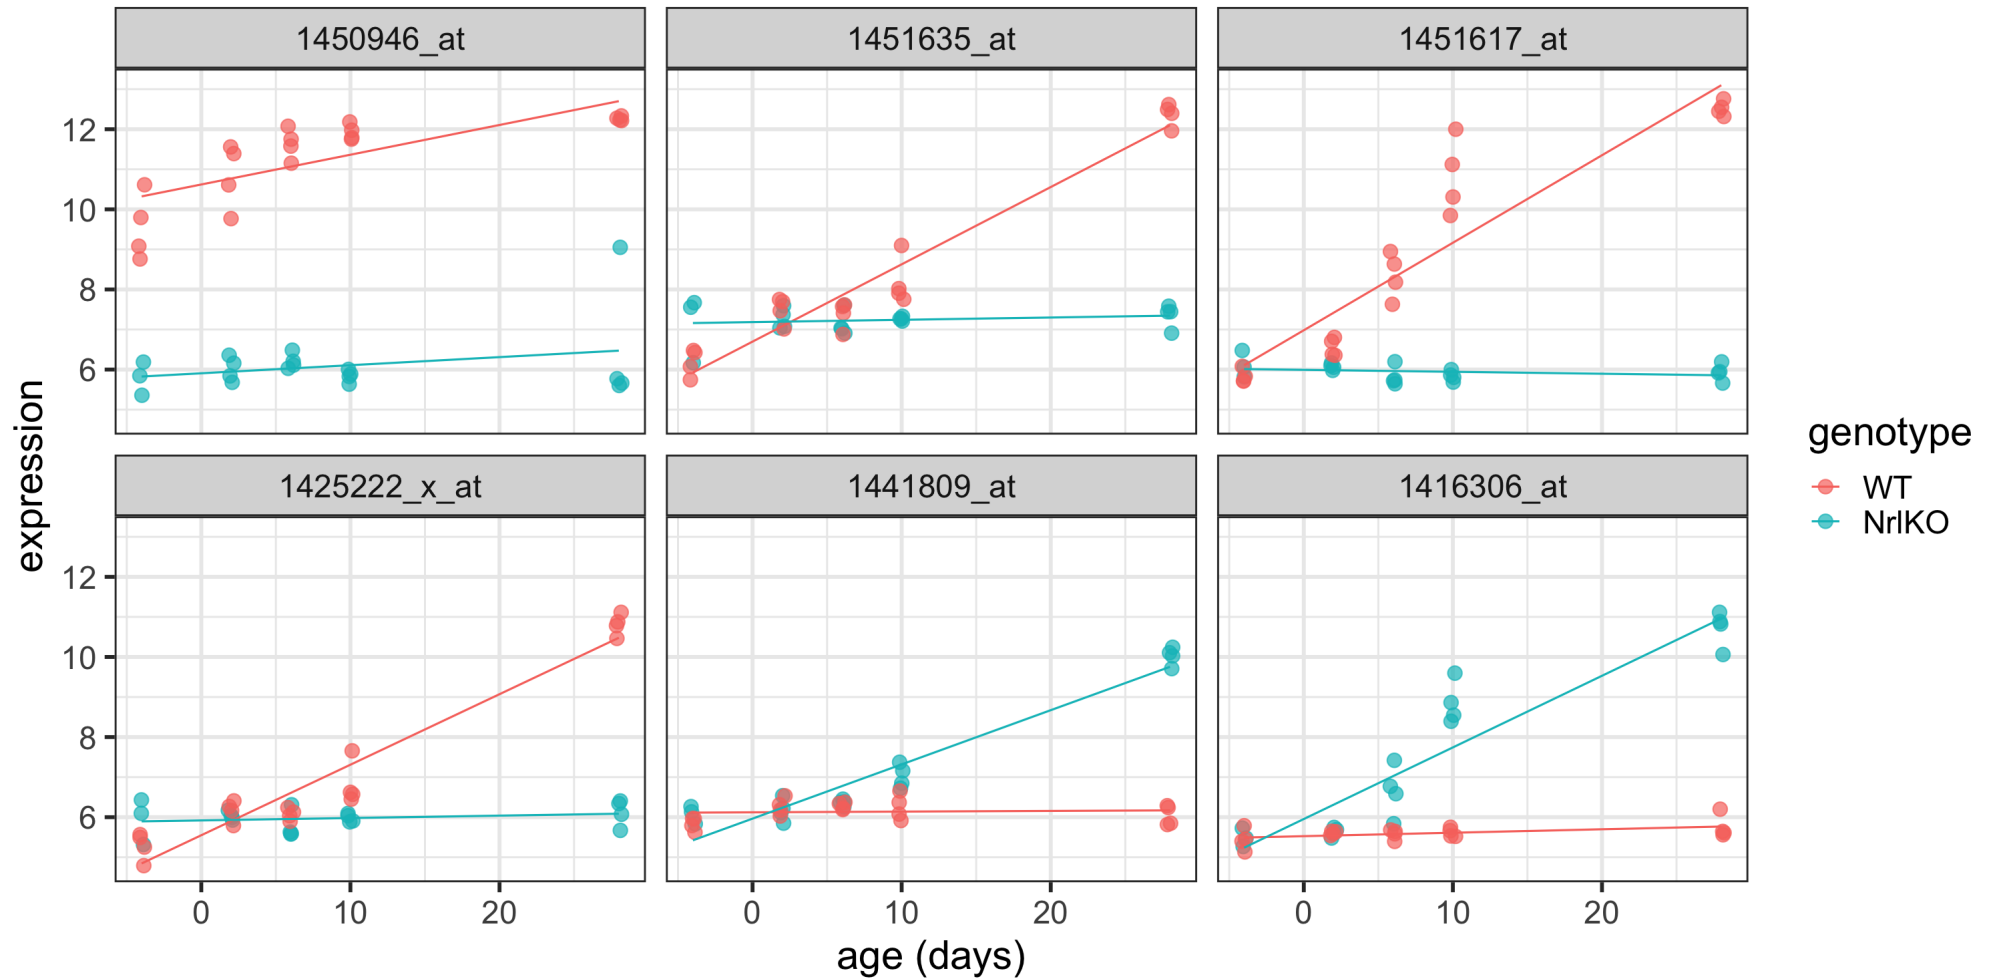
# topTable in action: any effect of genotype

```
topTable(ebFit, coef = c("genotypeNr1K0", "genotypeNr1K0:age"))[, -3]
```

##	genotypeNr1K0	genotypeNr1K0.age	F	P.Value	adj.P.Val
## 1450946_at	-4.7134891	-0.05411298	246.5762	2.403638e-22	1.084065e-17
## 1451635_at	0.4868705	-0.18718005	129.8112	1.307185e-17	2.180049e-13
## 1451617_at	-0.9885684	-0.22329233	128.9935	1.450111e-17	2.180049e-13
## 1425222_x_at	0.3698846	-0.17009077	120.0041	4.719613e-17	5.321482e-13
## 1441809_at	-0.1576064	0.13363563	117.5571	6.594308e-17	5.948198e-13
## 1416306_at	0.4278776	0.17020589	113.2125	1.212332e-16	9.112895e-13
## 1439143_at	0.6570269	0.06170447	111.2102	1.616023e-16	1.001158e-12
## 1449526_a_at	0.9935801	0.04232958	110.5598	1.775851e-16	1.001158e-12
## 1452243_at	0.1639373	-0.15313265	106.2784	3.344140e-16	1.675823e-12
## 1451763_at	-0.6819899	-0.14010359	102.7010	5.771358e-16	2.442992e-12

- `topTable(ebFit, coef = c(2,4))` is equivalent here, but much less informative!!
- What is the null hypothesis here?
  - $H_0 : \tau_{KO} = \tau_{KO:Age} = 0$
  - this finds genes where any (additive/interaction) effect of genotype is significant

# Plotting the top 6 genes for any effect of genotype




# Comparison of $s_g^2$ and $\tilde{s}_g^2$ (shrinkage!)

For **small** variances,  
limma \_\_\_\_\_ the  
estimates (fill in the  
blank):

(a) INCREASES

(b) DECREASES




# Comparison of $s_g^2$ and $\tilde{s}_g^2$ (shrinkage!)

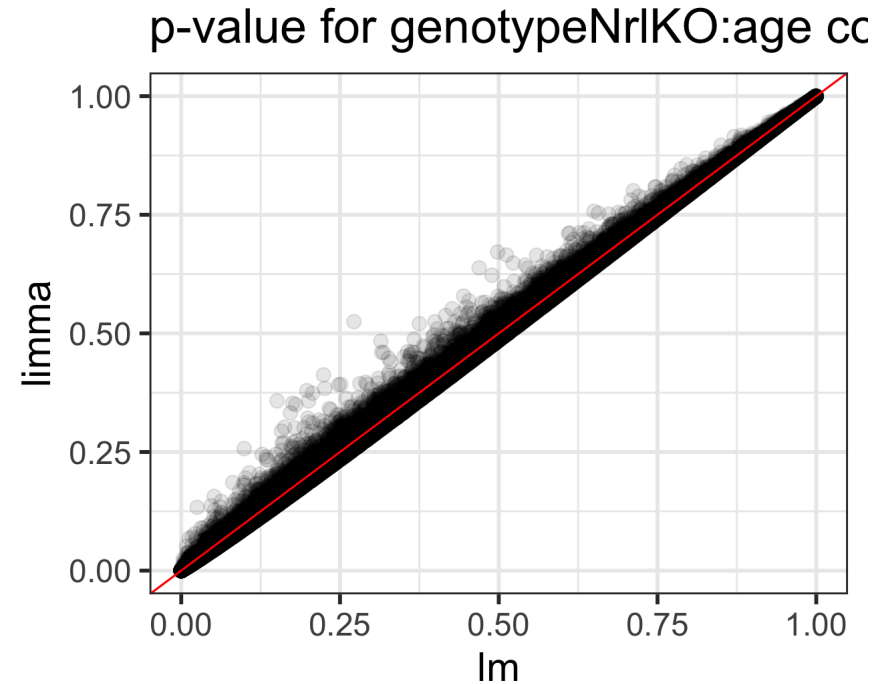
For **large** variances,  
limma \_\_\_\_\_ the  
estimates (fill in the  
blank):

(a) INCREASES

(b) DECREASES



# Comparison of interaction coefficient p-values



- 17479 genes where limma p-value is *larger* than lm
- 27622 genes where limma p-value is *smaller* than lm

# Multiple testing

# Error rates

Actual Situation “Truth”		
Decision	$H_0$ True	$H_0$ False
Do Not Reject $H_0$	Correct Decision $1-\alpha$	Incorrect Decision Type II Error $\beta$
Reject $H_0$	Incorrect Decision Type I Error $\alpha$	Correct Decision $1-\beta$

$$\alpha = P(\text{Type I Error}), \beta = P(\text{Type II Error}), \text{Power} = 1 - \beta$$


# Type I Error rate for $m$ tests

- $P(\text{incorrect decision}|H_0) = \alpha$ 
  - let  $\alpha = 0.05$
- $P(\text{correct decision}|H_0) = 1 - \alpha = 0.95$



# Type I Error rate for $m$ tests

- $P(\text{incorrect decision}|H_0) = \alpha$ 
  - let  $\alpha = 0.05$
- $P(\text{correct decision}|H_0) = 1 - \alpha = 0.95$
- $P(\text{correct decision on } m \text{ tests}|H_0) =$   
 $(1 - \alpha)^m = 0.95^m$
- $P(\text{at least one incorrect decision on } m \text{ tests}|H_0) =$   
 $1 - (1 - \alpha)^m =$   
 $1 - 0.95^m = \alpha_{FWER}$

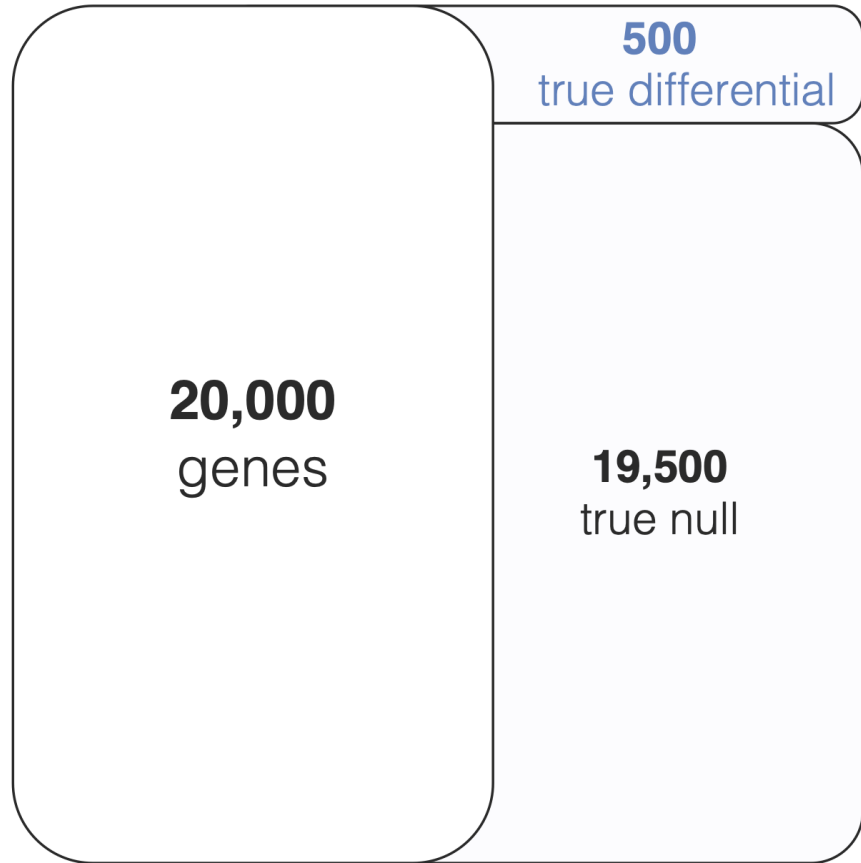


# Multiple comparisons and error rates

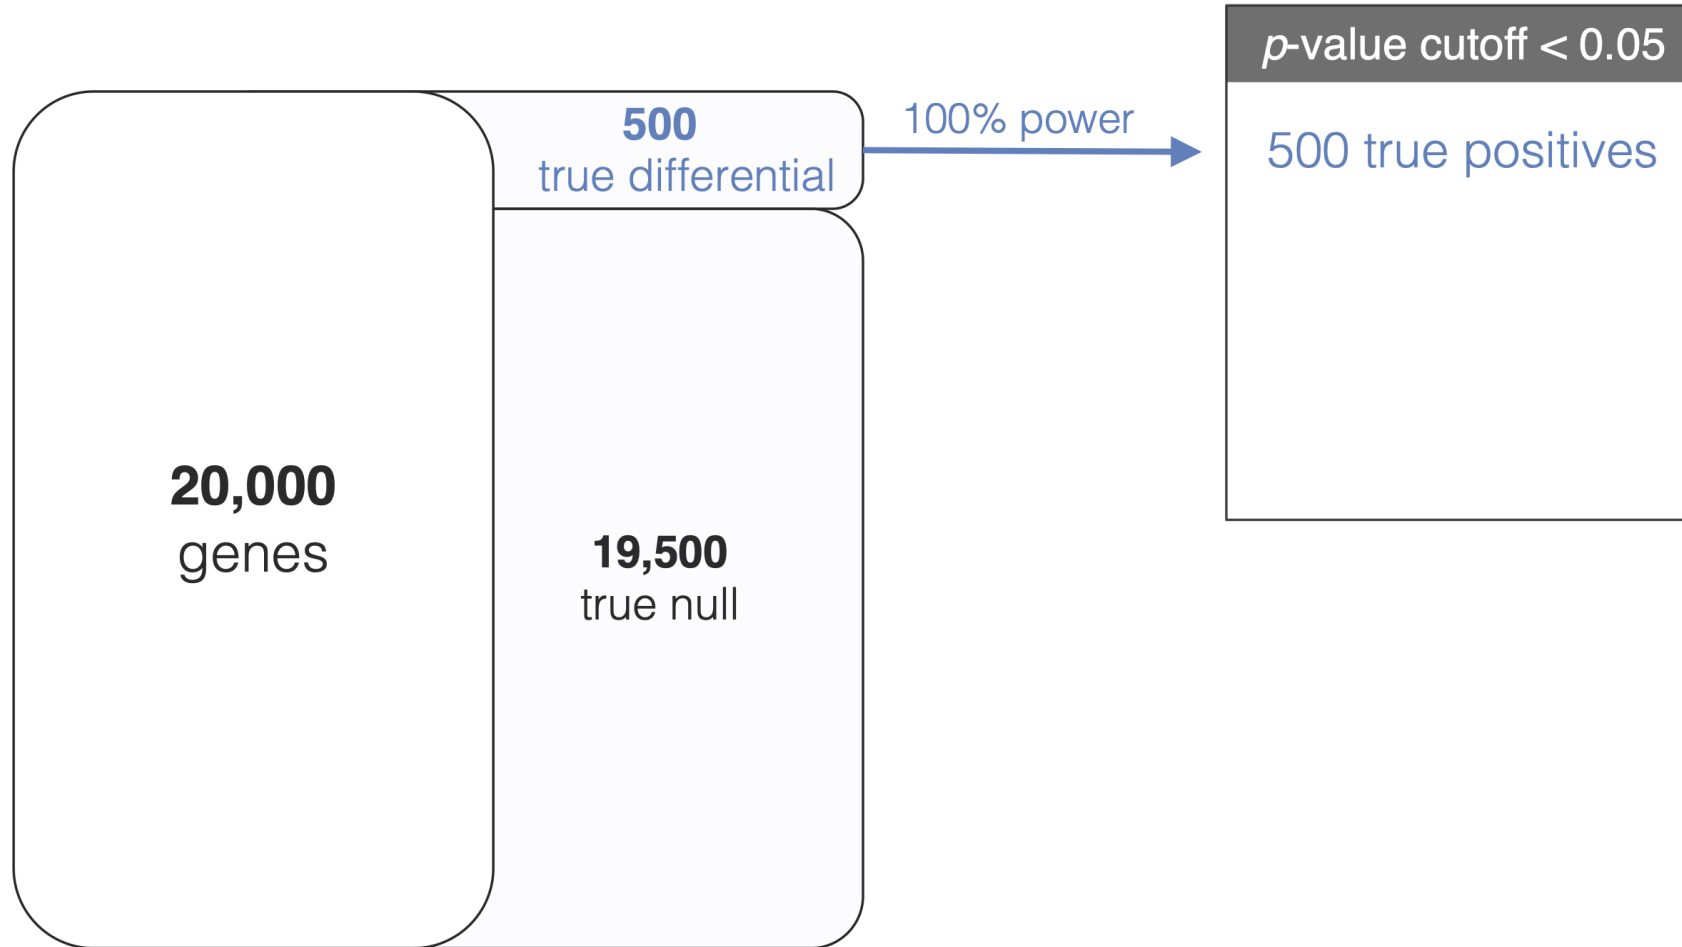


**20,000**  
genes

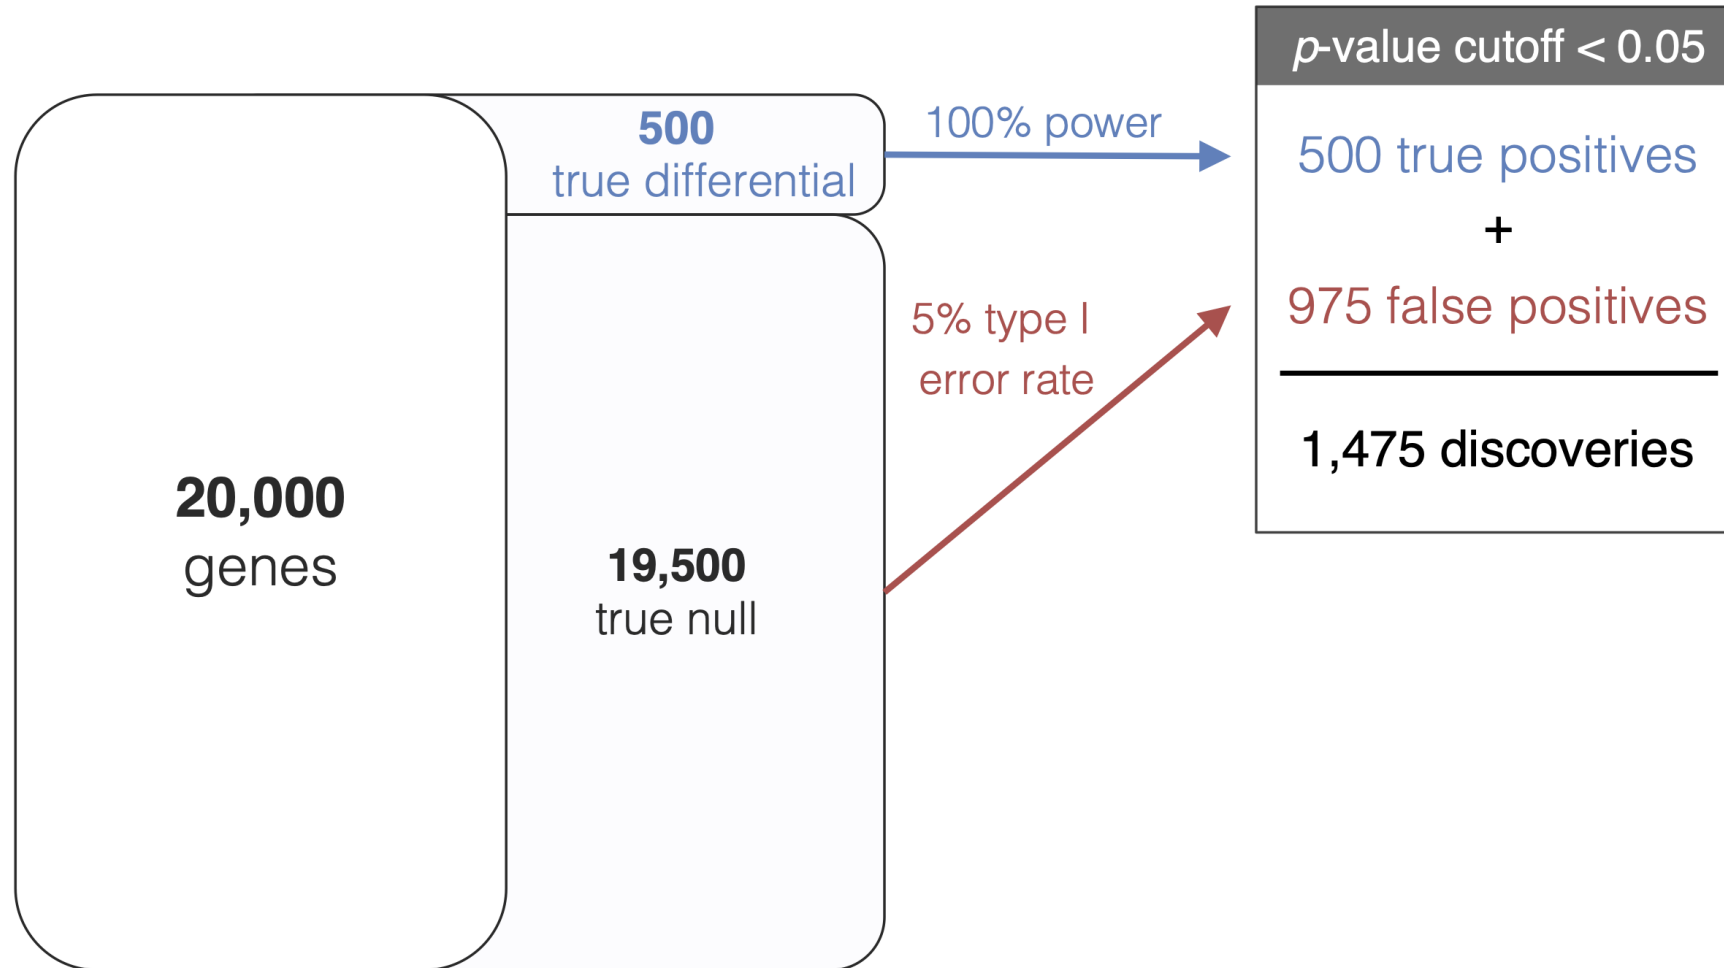
# Multiple comparisons and error rates



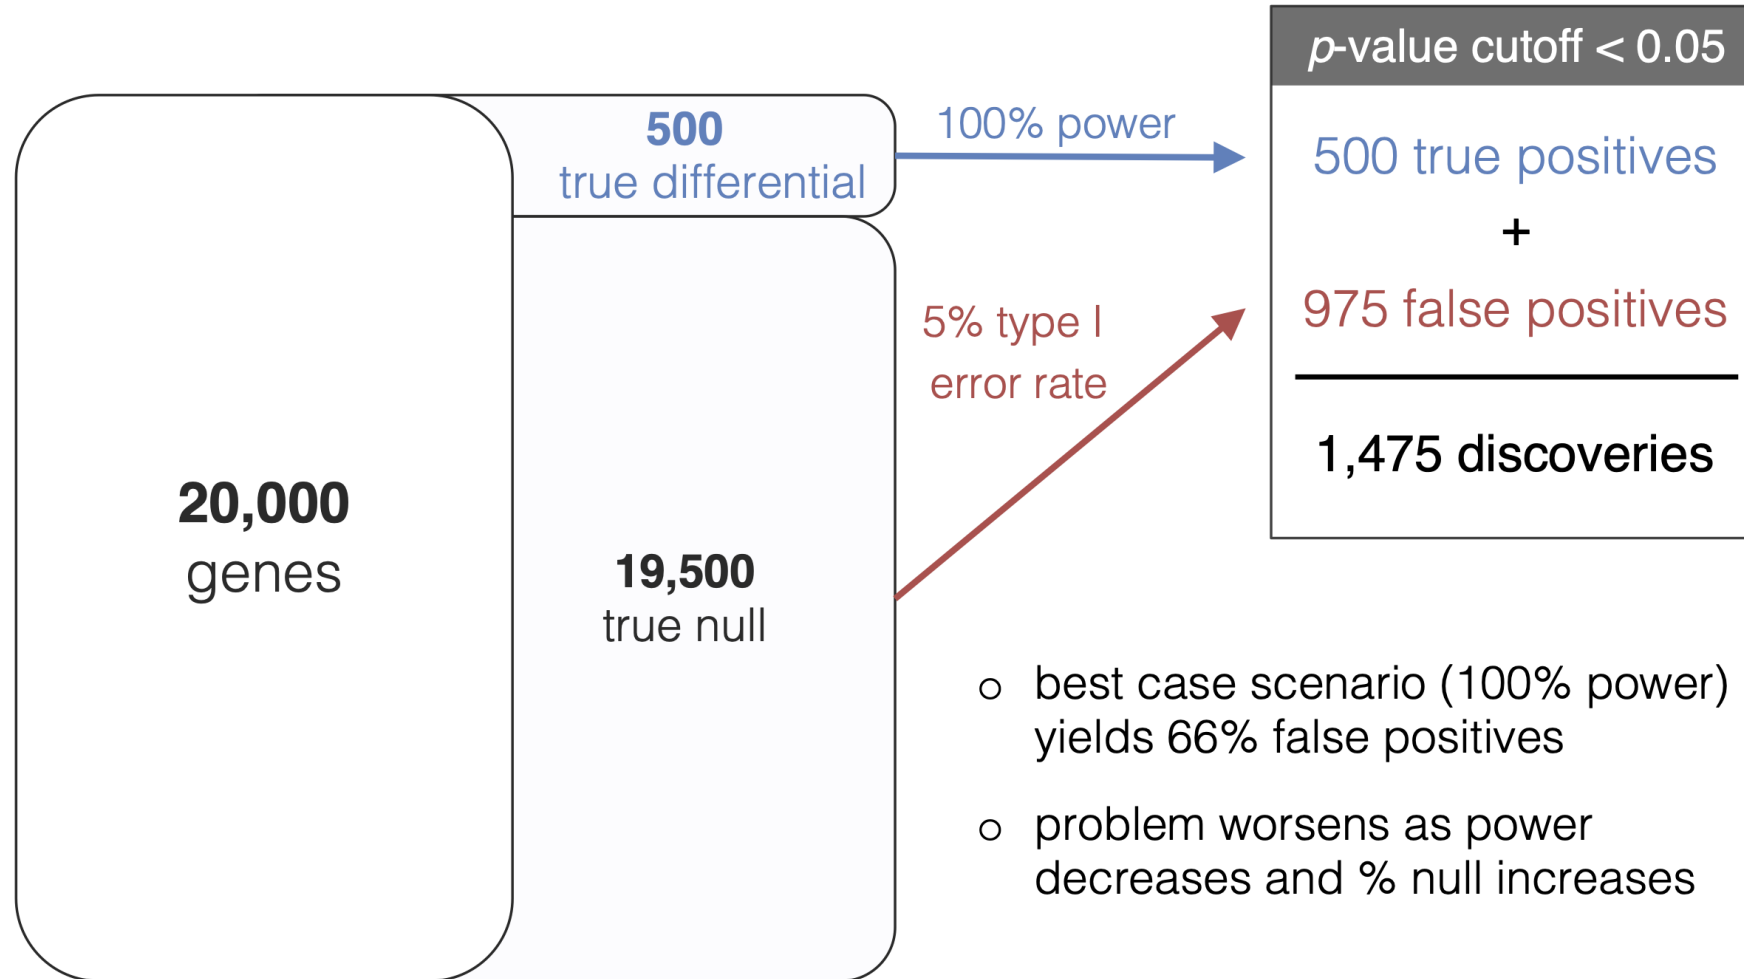
# Multiple comparisons and error rates



# Multiple comparisons and error rates



# Multiple comparisons and error rates



# Family-Wise Error Rate (FWER)

- **FWER** is the probability of making at least one error when testing  $m$  tests
- **Control the FWER**: limit the probability of making at least one incorrect decision
- One example: the **Bonferroni** correction for  $\alpha = 0.05$ :

If  $P(\text{at least one error on } m \text{ tests}) < \alpha$

$$\Rightarrow P(\text{at least one error on } m \text{ tests}) < \sum_{i=1}^m P(\text{error on test } i)$$

$$\sum_{i=1}^m P(\text{error on test } i) = m\alpha_{Bon}$$

$$\alpha_{Bon} = \frac{\alpha}{m} = \frac{0.05}{m}$$

# Bonferroni correction: controlling the FWER

Can think of controlling the probability of at least one false positive in two ways:

1. **Adjust the p-values; keep same  $\alpha$ :**

$$p_{Bon,i} = mp_i \text{ (more technically correct: } p_{Bon,i} = \min(mp_i, 1))$$

Then, threshold  $p_{Bon,i}$  at  $\alpha$

2. **Adjust the  $\alpha$  threshold; keep same p-values:**

$$\alpha_{Bon} = \frac{\alpha}{m}$$

Then, threshold  $p_i$  at  $\alpha_{Bon}$



# Multiple test correction is an active area of statistical research

- Bonferroni correction is very **conservative** (i.e. controls the FWER even lower than  $\alpha$  in many settings)
- Several other options are better
- For example, the Holm procedure: multiplier for p-value correction is not the same for all genes; more powerful

$$p_{Holm,1} = mp_1$$

$$p_{Holm,2} = (m - 1)p_2$$

$$p_{Holm,3} = (m - 2)p_3$$

$$\vdots$$

$$\Rightarrow FWER \leq \alpha$$

# How practical is the FWER in high-throughput biology?

- Why do we care so much about making one single error??



- One extreme way to ensure no Type I errors: reject no hypotheses! 😊
  - However, then our power is zero... 😭
- **Being overly strict about Type I error leads to greater Type II error (loss of power)**

**Radical idea: it's OK to make multiple mistakes, as long as you also have some true positives!**

# Enter: the False Discovery Rate (FDR)

*J. R. Statist. Soc. B* (1995)  
57, No. 1, pp. 289–300

## **Controlling the False Discovery Rate: a Practical and Powerful Approach to Multiple Testing**

By YOAV BENJAMINI† and YOSEF HOCHBERG

*Tel Aviv University, Israel*

[Received January 1993. Revised March 1994]

### SUMMARY

The common approach to the multiplicity problem calls for controlling the familywise error rate (FWER). This approach, though, has faults, and we point out a few. A different approach to problems of multiple significance testing is presented. It calls for controlling the expected proportion of falsely rejected hypotheses – the false discovery rate. This error rate is equivalent to the FWER when all hypotheses are true but is smaller otherwise. Therefore, in problems where the control of the false discovery rate rather than that of the FWER is desired, there is potential for a gain in power. A simple sequential Bonferroni-type procedure is proved to control the false discovery rate for independent test statistics, and a simulation study shows that the gain in power is substantial. The use of the new procedure and the appropriateness of the criterion are illustrated with examples.

Benjamini Y, Hochberg Y. "Controlling the false discovery rate: a practical and powerful approach to multiple testing." *Journal of the Royal statistical society: series B (Methodological)*. 1995 Jan;57(1):289-300.

Over 60K citations!!

# False Discovery Rate

	Null True	Alternative True	Total
Not Called Significant	$U$	$T$	$m - R$
Called Significant	$V$	$S$	$R$
	$m_0$	$m - m_0$	$m$

$V$  = # Type I errors [false positives]

FDR is designed to control the expected proportion of false positives ( $V$ ) among all hypotheses where the null has been rejected ( $R$ )

# False Discovery Rate

	Null True	Alternative True	Total
Not Called Significant	$U$	$T$	$m - R$
Called Significant	$V$	$S$	$R$
	$m_0$	$m - m_0$	$m$

$V$  = # Type I errors [false positives]

$$FDR = E\left[\frac{V}{R}\right]$$

# FDR vs FPR vs FWER

- **False Discovery Rate (FDR)** is the rate that significant features ( $R$ ) are truly null

$$FDR = E\left[\frac{V}{R}\right]$$

- **False Positive Rate (FPR)** is the rate that truly null features ( $m_0$ ) are called significant

$$FPR = E\left[\frac{V}{m_0}\right]$$

- **Family-Wise Error Rate (FWER)** is the probability that the number of truly null features rejected ( $V$ ) is at least 1

$$FWER = P(V \geq 1)$$

# Benjamini-Hochberg FDR (BH procedure)

- Proposed the idea of controlling FDR instead of FWER
- Proposed a procedure for doing so
  - note that we know  $R$ , but we don't know  $V$
- Procedure: control FDR at level  $q$ 
  1. order the raw p-values  $p_1 \leq p_2 \leq \dots \leq p_m$
  2. find test with highest rank  $j$  such that  $p_j < \frac{jq}{m}$
  3. declare all smaller ranks up to  $j$  significant



# Controlling FDR at level $q = 0.05$

Rank ( $j$ )	P-value
1	0.0008
2	0.009
3	0.127
4	0.205
5	0.396
6	0.450
7	0.641
8	0.781
9	0.900
10	0.993

# Controlling FDR at level $q = 0.05$

Rank ( $j$ )	P-value	$(j/m)*q$
1	0.0008	0.005
2	0.009	0.010
3	0.127	0.015
4	0.205	0.020
5	0.396	0.025
6	0.450	0.030
7	0.641	0.035
8	0.781	0.040
9	0.900	0.045
10	0.993	0.050

# Controlling FDR at level $q = 0.05$

Rank ( $j$ )	P-value	$(j/m)*q$	Reject $H_0$ ?
1	0.0008	0.005	✓
2	0.009	0.010	✓
3	0.127	0.015	
4	0.205	0.020	
5	0.396	0.025	
6	0.450	0.030	
7	0.641	0.035	
8	0.781	0.040	
9	0.900	0.045	
10	0.993	0.050	

# Controlling FDR at level $q = 0.05$

Rank ( $j$ )	P-value	$(j/m)*q$	Reject $H_0$ ?	$FWER_{Bon} < 0.05$ ?
1	0.0008	0.005	✓	✓
2	0.009	0.010	✓	
3	0.127	0.015		
4	0.205	0.020		
5	0.396	0.025		
6	0.450	0.030		
7	0.641	0.035		
8	0.781	0.040		
9	0.900	0.045		
10	0.993	0.050		

Where  $\alpha_{Bon} = 0.05/10 = 0.005$

# BH FDR values given in `limma` by default

```
topTable(ebFit, coef = "genotypeNr1K0")
```

##		logFC	AveExpr	t	P.Value	adj.P.Val	B
##	1450946_at	-4.7134891	8.733103	-15.583683	5.390931e-18	2.431364e-13	28.40938
##	1442070_at	-0.9669400	8.268122	-9.787407	6.896220e-12	1.555132e-07	16.46086
##	1443457_at	-1.1530057	6.049428	-9.091412	4.947029e-11	6.029428e-07	14.68843
##	1422679_s_at	-2.1187394	9.495435	-9.002625	6.388954e-11	6.029428e-07	14.45703
##	1431708_a_at	-2.1610625	8.591450	-8.967950	7.062007e-11	6.029428e-07	14.36634
##	1433050_at	-1.5083887	8.468891	-8.923932	8.021234e-11	6.029428e-07	14.25096
##	1426288_at	-4.2100657	9.323796	-8.772209	1.246491e-10	8.031144e-07	13.85106
##	1418108_at	0.8662091	8.260647	8.537984	2.475294e-10	1.395478e-06	13.22713
##	1450770_at	1.3103576	7.357033	8.128847	8.332369e-10	3.747976e-06	12.11870
##	1457802_at	-0.8362030	7.778189	-8.117112	8.629984e-10	3.747976e-06	12.08657

Or, obtain them yourself for any vector of p-values `p` with `p.adjust(p, method="BH")`

# Other ways to control FDR

- BH is just one (the first) method to control FDR
- Since the publication of the BH method, other methods have been proposed
- One of the most popular is Storey's  $q$ -value

## Statistical significance for genomewide studies

John D. Storey<sup>\*†</sup> and Robert Tibshirani<sup>†</sup>

<sup>\*</sup>Department of Biostatistics, University of Washington, Seattle, WA 98195; and <sup>†</sup>Departments of Health Research and Policy and Statistics, Stanford University, Stanford, CA 94305

Edited by Philip P. Green, University of Washington School of Medicine, Seattle, WA, and approved May 30, 2003 (received for review January 28, 2003)

With the increase in genomewide experiments and the sequencing of multiple genomes, the analysis of large data sets has become commonplace in biology. It is often the case that thousands of features in a genomewide data set are tested against some null hypothesis, where a number of features are expected to be significant. Here we propose an approach to measuring statistical significance in these genomewide studies based on the concept of the false discovery rate. This approach offers a sensible balance between the number of true and false positives that is automatically calibrated and easily interpreted. In doing so, a measure of statistical significance called the  $q$  value is associated with each tested feature. The  $q$  value is similar to the well known  $p$  value, except it is a measure of significance in terms of the false discovery rate rather than the false positive rate. Our approach avoids a flood of false positive results, while offering a more liberal criterion than what has been used in genome scans for linkage.

false discovery rates | genomics | multiple hypothesis testing |  $q$  values

to the method in ref. 5 under certain assumptions. Also, ideas similar to FDRs have appeared in the genetics literature (1, 13).

Similarly to the  $p$  value, the  $q$  value gives each feature its own individual measure of significance. Whereas the  $p$  value is a measure of significance in terms of the false positive rate, the  $q$  value is a measure in terms of the FDR. The false positive rate and FDR are often mistakenly equated, but their difference is actually very important. Given a rule for calling features significant, the false positive rate is the rate that truly null features are called significant. The FDR is the rate that significant features are truly null. For example, a false positive rate of 5% means that on average 5% of the truly null features in the study will be called significant. A FDR of 5% means that among all features called significant, 5% of these are truly null on average.

The  $q$  value provides a measure of each feature's significance, automatically taking into account the fact that thousands are simultaneously being tested. Suppose that features with  $q$  values  $\leq 5\%$  are called significant in some genomewide test of significance. This results in a FDR of 5% among the significant features. A

- `qvalue` package implementation: provides adjusted  $p$ -values

# Storey's q-value vs BH (Conceptual)

- Just like BH, is focused on the proportion of discoveries that are false positives
- *Conceptual* difference between BH and Storey's q-value is:
  - BH **controls** the FDR
  - q-values give an unbiased **estimate** of the FDR (will control the FDR on eaverage)

# Storey's q-value vs BH (Mathematical)

- Mathematically, the difference between the two is in how  $m_0$  is estimated
  - Or equivalently, how  $\pi_0 = \frac{m_0}{m}$  is estimated (since  $m$  is known)
  - $\pi_0$  represents the proportion of tests that are truly null
- q-value:

$$\hat{q}(p_i) = \min_i \left( \frac{\hat{\pi}_0 m}{\text{rank}(p_i)} p_i, 1 \right)$$

- q-value and BH-adjusted p-values are equivalent when  $\pi_0 = 1$

$$\hat{p}_{BH}(p_i) = \min_i \left( \frac{m}{\text{rank}(p_i)} p_i, 1 \right)$$

(BH conservatively assumes that  $\pi_0 = 1$ )



# BH vs q-value in our example

Rank ( $j$ )	P-value	$\hat{p}_{BH}(p_i)$	$\hat{q}(p_i)$
1	0.0008	0.008	0.008
2	0.009	0.045	0.045
3	0.127	0.423	0.423
4	0.205	0.513	0.513
5	0.396	0.792	0.750
6	0.450	0.750	0.750
7	0.641	0.916	0.916
8	0.781	0.976	0.976
9	0.900	1.000	0.993
10	0.993	0.993	0.993

# Compounding issues of multiple comparisons


- What if you're not only testing 45K genes, but also multiple tests per gene (e.g. multiple contrasts, such as several two-group comparisons)?
- Classical procedures for adjustment:
  - Tukey multiple comparison procedure
  - Scheffe multiple comparison procedure
  - Bonferroni or Holm FWER correction
- In our setting, we can also apply BH to all p-values globally
  - `limma::decideTests(pvals, method="global")` for a matrix of p-values or eBayes output (e.g. rows = genes, columns = contrasts)
  - p-values are combined, adjusted globally, then separated back out and sorted

# Assumptions about p-values

- Implicit assumption for all multiple testing correction methods:


**p-value distribution is "well-behaved"**

- What does this mean?
  - primarily, that the distribution of p-values under the null is **uniform**



# p-value distributions

Spike of small p-values indicates non-null tests:



Great primer on how things can go wrong: <http://varianceexplained.org/statistics/interpreting-pvalue-histogram/>

# What if p-values are *poorly behaved*?

- FDR estimates can be invalid (assumptions are violated)
- Solution: compute p-values "empirically" using resampling/permutation/bootstrap techniques
- **Bootstrap:** take repeated random samples with replacement from your data and compute statistic; repeat many times and use bootstrap statistics as your sampling distribution rather than a t, Normal, F,  $\chi^2$ , etc
- **Permutation:** construct a simulated version of your dataset that satisfies the null hypothesis and compute statistic (e.g. shuffle group labels for a two-group comparison); repeat many times and use permutation statistics as your sampling distribution rather than a t, Normal, F,  $\chi^2$ , etc
- Downside: often computationally intensive for genomics