

# Statistical Methods for High Dimensional Biology

## Supervised learning I: Classification

Keegan Korthauer

10 March 2021

with slide contributions from Sara Mostafavi, Gabriela Cohen-Freue, and Kevin Murphy

# Announcements

- Project presentation dates posted today
- Project progress reports due Friday 19 March
- Next week's lectures (15 and 17 March) will both be **synchronous**
  - Guest lecturer Paul Pavlidis will talk about two topics that will be very relevant for many projects: Gene set enrichment analysis, and Gene networks and function prediction

# Learning objectives

- Explain the purpose of **supervised learning** and how it differs from **unsupervised learning**
- Connect commonly used terms from statistics and machine learning
- Explain the goals of **classification**
- Understand the main ideas behind the mathematical frameworks such as **Naïve Bayes**, **Linear Discriminant Analysis**, and **K-Nearest Neighbors classification**

# Supervised learning

A procedure or algorithm which uses a set of **inputs** (measured or preset variables) to predict the values of one or more **outputs** (variables which are influenced in some way by the inputs)

- This definition uses the language/terminology from the field of *machine learning*
- In statistical terminology:
  - we would say *predictor* or *independent* variables in place of **inputs**
  - we would say *response* or *dependent* variables in place of **outputs**

# Examples in genomics

# Machine learning vs classical statistics

## Machine Learning

- Large number of variables (have no idea which are useful)
- Model complex, non-linear relationships
- Flexibility about defining a classifier: “loss/error minimization view” vs “generative” view
- Invent scalable algorithms that can solve parameters for very large models

## Classical Statistics

- Handful of variables
- Typically assume linear relationships
- Typically think in terms of a “generative” model; has theoretical justification
- Thorough analysis/theory for models with less than a dozen parameters

# ML vs Statistics terminology

## Machine Learning

- Labels / 'class' labels
- Examples
- Features
- Learning
- Weights / feature importance
- Generalization

## Classical Statistics

- Response / outcome
- Data points
- Covariates / variables
- Estimation / fitting
- Parameters
- Test set performance

# Example: Predict phenotype from gene expression

1. Measure gene expression data relevant for the outcome you would like to predict (e.g., disease status or severity): **training data**
2. Formulate (i.e., write down) a **model** that relates the gene expression measurements (i.e. features/attributes) to the outcome
3. Fit/estimate model **parameters** based on data to fully specify the model
4. Apply the model to new data, where you don't have (or at least don't *use*) information about outcome/response to make a prediction



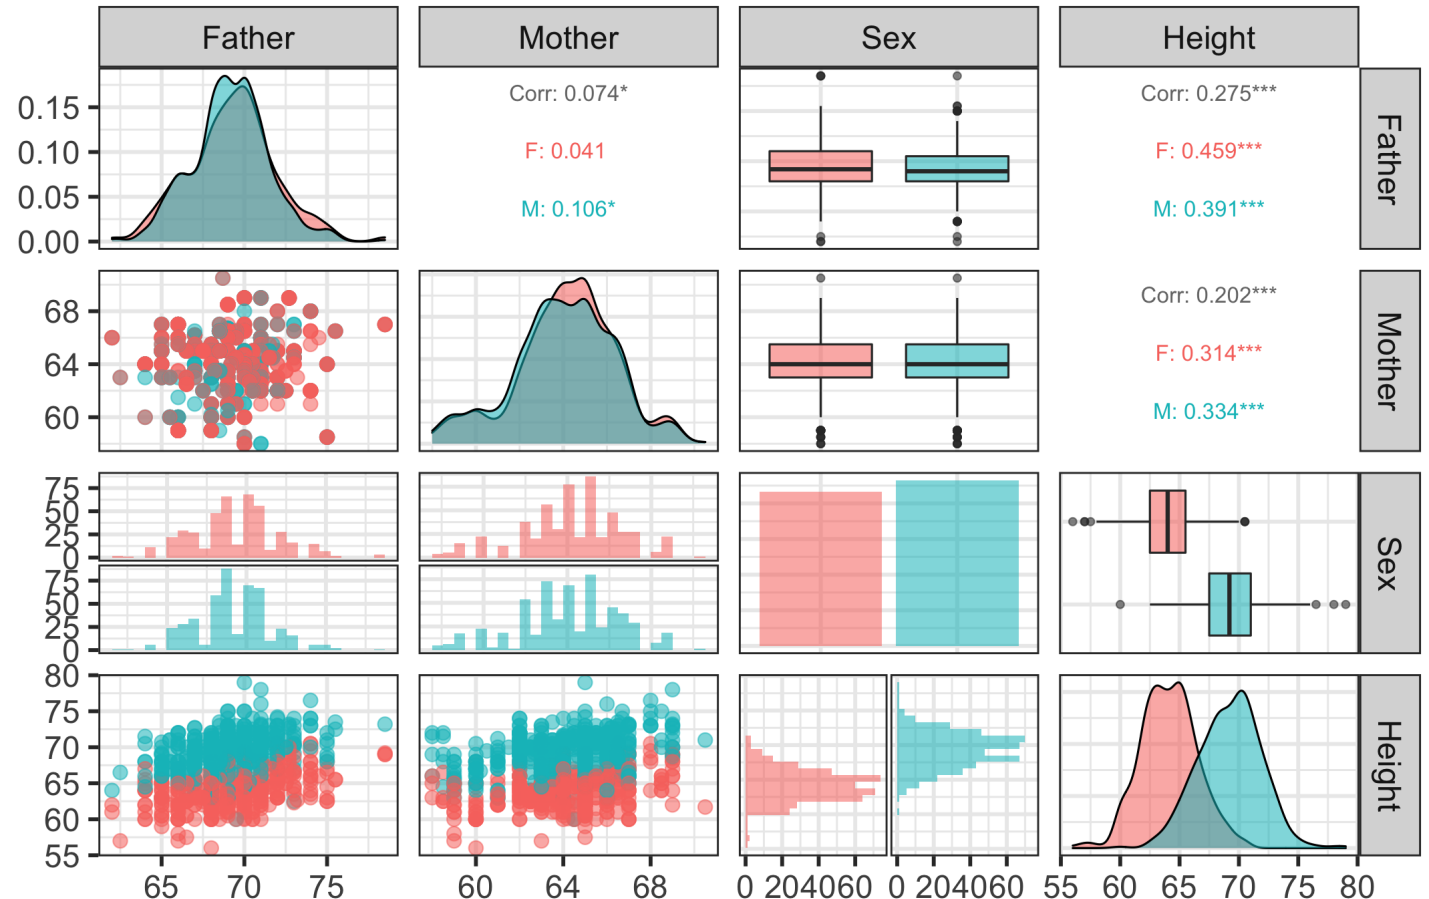
# Galton's Height Data: predict the future (adult) height of a child

Data source

# 1. Gather training data

$\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$

- **input:**  $x_i$ , feature vector (Father height, mother height, sex)
- **output:**  $y_i$ , response (child height)



## 2. Formulate model relating input and output

- Write down a model that links the output variable to some function of the input variable(s)

$$y_i = f(\mathbf{x}_i) + \epsilon$$

- For example, let's say child height is linearly related to the additive effects of parental mean height and sex

$$y_i = \alpha + \beta_1 \left( \frac{x_{father,i} + x_{mother,i}}{2} \right) + \beta_2 x_{male,i} + \epsilon$$

- $x_{father,i}$  and  $x_{mother,i}$  are the heights of the father and mother of child  $i$ , and  $x_{male,i}$  is an indicator variable that the  $i^{th}$  child's sex is male
- $\alpha$ ,  $\beta_1$ , and  $\beta_2$  are model parameters

### 3. Fit the model to training data

- Let  $\bar{x}_{parental,i}$  represent parental mean height
- How can we fit this model to minimize *error* on the training data?

$$y_i = \alpha + \beta_1 \bar{x}_{parental,i} + \beta_2 x_{male,i} + \epsilon_i, \text{ where } \epsilon_i \sim N(0, \sigma^2)$$

- e.g. find  $\alpha$ ,  $\beta_1$ , and  $\beta_2$  such that the objective function (sum of squared errors) is minimized

$$\sum_{i=1}^n (y_i - \alpha - \beta_1 \bar{x}_{parental,i} - \beta_2 x_{male,i})^2$$

### 3. Fit the model to training data

- Let  $\bar{x}_{parental,i}$  represent parental mean height
- How can we fit this model to minimize *error* on the training data?

$$y_i = \alpha + \beta_1 \bar{x}_{parental,i} + \beta_2 x_{male,i} + \epsilon_i, \text{ where } \epsilon_i \sim N(0, \sigma^2)$$

- e.g. find  $\alpha$ ,  $\beta_1$ , and  $\beta_2$  such that the objective function (sum of squared errors) is minimized

$$\sum_{i=1}^n (y_i - \alpha - \beta_1 \bar{x}_{parental,i} - \beta_2 x_{male,i})^2$$

Linear regression!

# Aside: generative model

(Mathematically equivalent) probabilistic formulation of linear regression:

$$y_i | \bar{x}_{parental,i}, x_{male,i} \sim N(\alpha + \beta_1 \bar{x}_{parental,i} + \beta_2 x_{male,i}, \sigma^2)$$

$$p(y_1, y_2, \dots, y_n | \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) = \prod_{i=1}^n f_N(y_i | \alpha + \beta_1 \bar{x}_{parental,i} + \beta_2 x_{male,i}, \sigma^2)$$

Where  $f_N(y | \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}(\frac{y-\mu}{\sigma})^2}$  is the normal probability density function

## 4. Apply model to predict on new data

- Suppose we collect a new dataset of parental mean heights  $\bar{x}_{parental,i}^*$  and child's sex  $x_{male,i}^*$  for an independent set of  $m$  children  $i = 1, \dots, m$
- We would like to predict each child's eventual adult height  $\hat{y}_i^*$  based their sex and their parents' mean height
- **How?** Use the model parameters estimated from the training data  $(\hat{\alpha}, \hat{\beta}_1, \hat{\beta}_2)$  and plug in our predictor variables  $(\bar{x}_{parental,i}^*, x_{male,i}^*)$

$$\hat{y}_i^* = \hat{\alpha} + \hat{\beta}_1 \bar{x}_{parental,i}^* + \hat{\beta}_2 x_{male,i}^*$$

# Supervised learning

## Regression

- continuous outcome

## Classification

- binary outcome
- categorical outcome



# The classification problem

Training data:

$$\{(\mathbf{x}_1, c_1), (\mathbf{x}_2, c_2), \dots, (\mathbf{x}_n, c_n)\}$$

Instead of a *continuous* outcome/response (e.g. height), we now have discrete **class labels**  $c_i \in \{1, \dots, K\}$

# Classifiers


A **classifier** is a function  $f$  that maps input feature vectors  $\mathbf{x}_i = \{x_{i1}, x_{i2}, \dots, x_{ip}\}$  to output class labels  $c_i \in \{1, \dots, K\}$

- we assume that class  $i$  labels are *unordered* and *mutually exclusive*
- Let  $\mathcal{X}$  be the feature space of all possible values of  $\mathbf{x}_i$ 
  - this space could consist of continuous, discrete values or a mixture of the two (e.g.  $\mathcal{X} = \{0, 1\}^p$  or  $\mathcal{X} = \mathbb{R}^p$ )
- **Goal:** to learn a function  $f$  that maps feature vectors to labels, based on labeled training set:  $\{(\mathbf{x}_1, c_1), (\mathbf{x}_2, c_2), \dots, (\mathbf{x}_n, c_n)\}$

$$f(\mathbf{x}_i) = c_i$$

# Example classification task


Partition the space of input data so that we minimize the number of “miss-classified” objects/points



# Example classification task

Partition the space of input data so that we minimize the number of “miss-classified” objects/points


Our classification boundary might be more or less complex (e.g. nonlinear vs linear)




# Example classification task

Partition the space of input data so that we minimize the number of “miss-classified” objects/points

Our classification boundary might be more or less complex (e.g. nonlinear vs linear)



# More than 2 groups



# Defining the cost of misclassification

To train a model, we need a **loss / error** function

The **cost** of a misclassification can be specified with a loss (or error) function

For example, if all misclassification errors are equally bad, the loss function for predicted class labels  $\hat{c}_i$  given true class labels  $c_i$  would be

$$L(c_i, \hat{c}_i) = \begin{cases} 0 & \text{if } c_i = \hat{c}_i \\ 1 & \text{otherwise} \end{cases}$$

# Minimizing the expected loss

**Goal:** predict the class that minimizes the *conditional expected loss*

- Expected conditional loss:  $\rho(\hat{c}(\mathbf{x}_i)) = \sum_{k=1}^K L(k, \hat{c}(\mathbf{x}_i))P(C = k|\mathbf{x}_i)$
- Simplest case of two classes:
  - Expected loss of predicting class 1 when label was 2:  $L(2, 1)p(2|\mathbf{x}_i)$
  - Expected loss of predicting class 2 when label was 1:  $L(1, 2)p(1|\mathbf{x}_i)$
  - Predict class 1 if:  $L(2, 1)p(2|\mathbf{x}_i) < L(1, 2)p(1|\mathbf{x}_i)$
- But how to get  $P(C = k|\mathbf{x}_i)$ ?



# Three main ways to solve classification problem

1. Learn a **generative model** (function) for the probability distribution of inputs for each class:  $p(\mathbf{x}_i | C = k)$ 
  - Then use **Bayes rule** and the marginal distribution  $p(C = k)$  (overall class prevalence) to predict  $p(C = k | \mathbf{x}_i)$
  - Recall Bayes rule:  $P(A|B) = \frac{P(B|A)P(A)}{P(B)}$
2. Learn a **discriminative model** for conditional probability distribution of each class  $p(C = k | \mathbf{x}_i)$ 
  - Do not consider the distribution of the predictors  $\mathbf{x}_i$
3. Learn a **non-parametric model**
  - e.g. a function that directly maps  $\mathbf{x}_i$  to its predicted class  $c$

# Example classification methods

## 1. Generative

- **Naïve Bayes**
- **Linear/Quadratic discriminant analysis**

## 2. Discriminative

- **Logistic regression**
- Support vector machines
- Decision trees (and Random Forest)
- Neural networks

## 3. Non-parametric

- **K-nearest neighbors**

# Generative model solution

Learn the following for each value of  $k \in \{1, \dots, K\}$ :

- **class-conditional density**  $p(\mathbf{x}_i | C = k)$
- **class priors** (overall class prevalence)  $p(C = k)$

Then apply Bayes rule to compute most likely class for each object/entity (*posterior*)

$$\begin{aligned} p(C = k | \mathbf{x}_i) &= \frac{p(\mathbf{x}_i | C = k)p(C = k)}{p(\mathbf{x}_i)} \\ &= \frac{p(\mathbf{x}_i | C = k)p(C = k)}{\sum_{j=1}^K p(\mathbf{x}_i | C = j)p(C = j)} \end{aligned}$$

# Naïve Bayes

- Most general of the **generative model** techniques
- Assumes features  $\{x_{i1}, x_{i2}, \dots, x_{ip}\}$  are **independent**
- Since features are independent, the conditional density of features given class can be written as the product of the individual feature conditional densities:  
$$p(\mathbf{x}_i | C = k) = \prod_{m=1}^p p(x_{im} | C = k)$$
- Useful only when the number of predictors is small (otherwise, hard to estimate all conditional distributions of features given class  $p(x_{im} | C = k)$ )

# Gaussian Naïve Bayes

Assume features within each class are *independently* normally (or Gaussian) distributed


$$\begin{aligned} p(\mathbf{x}_i | C = k) &= \prod_{m=1}^p p(x_{im} | C = k) \\ &= \prod_{m=1}^p f_N(x_{im} | \mu_{mk}, \sigma_{mk}) \end{aligned}$$

where  $f_N(x | \mu, \sigma)$  is the Normal probability density function with mean  $\mu$  and sd  $\sigma$





Estimate  $\boldsymbol{\mu}_k = \{\mu_{1k}, \dots, \mu_{pk}\}$ , and  $\boldsymbol{\sigma}_k = \{\sigma_{1k}, \dots, \sigma_{pk}\}$  using maximum likelihood estimation (MLE) on training data

# Example: Height data

Data/Example modified from: [Intro to Data Science by Irizarry](#)



# Class-conditional densities of features



# Class-conditional densities of features

```
(mles <- heights %>%  
  group_by(sex) %>%  
  summarize(mean_height = mean(height),  
            sd_height = sd(height),  
            mean_weight = mean(weight),  
            sd_weight = sd(weight)))
```

```
## # A tibble: 2 x 5  
##   sex      mean_height sd_height mean_weight sd_weight  
## * <fct>      <dbl>      <dbl>      <dbl>      <dbl>  
## 1 Female      64.9        3.76      137.       12.5  
## 2 Male       69.3        3.61      174.       16.9
```

e.g.  $x_{i,weight} | C = Female \sim N(137, 12.5)$



# Prediction using Gaussian Naïve Bayes

Goal: compute the **posterior probability** of class assignment for some new observation  $\mathbf{x}^*$

$$\begin{aligned} p(C = k | \mathbf{x}_i) &= \frac{p(\mathbf{x}_i | C = k) p(C = k)}{\sum_{j=1}^K p(\mathbf{x}_i | C = j) p(C = j)} \\ &= \frac{p(C = k) \prod_{m=1}^p p(x_{im} | C = k)}{\sum_{j=1}^K p(C = j) \prod_{m=1}^p p(x_{im} | C = j)} \\ &= \frac{p(C = k) \prod_{m=1}^p f_N(x_{im} | \mu_{mk}, \sigma_{mk})}{\sum_{j=1}^K p(C = j) \prod_{m=1}^p f_N(x_{im} | \mu_{mj}, \sigma_{mj})} \end{aligned}$$

# Prediction using Gaussian Naïve Bayes


In our two-class example with heights,

$$p(C = k | \mathbf{x}_i) = \frac{p(C = k) \prod_{m=1}^p f_N(x_{im} | \mu_{mk}, \sigma_{mk})}{p(C = 1) \prod_{m=1}^p f_N(x_{im} | \mu_{m1}, \sigma_{m1}) + p(C = 2) \prod_{m=1}^p f_N(x_{im} | \mu_{m2}, \sigma_{m2})}$$

- $p(C = 1)$  = overall proportion in class "Female" in training data (0.227)
- $p(C = 2)$  = overall proportion in class "Male" in training data (0.773)
- Male class height distribution:  $x_{i,height} | \mu_{height,2}, \sigma_{height,2} \sim N(69.3, 3.61)$
- Likewise we also have the female class height distribution, and the male and female class weight distributions (all Gaussian)

# Prediction using Gaussian Naïve Bayes

For example, let's say we have a new observation:  $\mathbf{x}^* = (x_{height}^* = 68, x_{weight}^* = 160)$



# Prediction using Gaussian Naïve Bayes

For example, if  $\mathbf{x}^* = (x_{height}^* = 68, x_{weight}^* = 160)$  we can compute  $p(C = 2|\mathbf{x}^*)$  (predicted probability that this sample is male):

```
new_height <- 68
new_weight <- 160
pM <- sum(heights == "Male")/nrow(heights)
mle_male <- mles %>% filter(sex == "Male")
mle_female <- mles %>% filter(sex == "Female")
num <- pM * (dnorm(new_height, mle_male$mean_height, mle_male$sd_height)*
             dnorm(new_weight, mle_male$mean_weight, mle_male$sd_weight))
den <- num +
  (1-pM) * (dnorm(new_height, mle_female$mean_height, mle_female$sd_height)*
            dnorm(new_weight, mle_female$mean_weight, mle_female$sd_weight))

num/den
```

```
## [1] 0.9288177
```

# Prediction using Gaussian Naïve Bayes

Prediction/fitting on entire training dataset (classification rule assigns each observation to the class with posterior probability  $> 0.5$ )

```
num <- pM * (dnorm(heights$height, mle_male$mean_height, mle_male$sd_height)*
             dnorm(heights$weight, mle_male$mean_weight, mle_male$sd_weight))
den <- num +
      (1-pM) * (dnorm(heights$height, mle_female$mean_height, mle_female$sd_height)*
                dnorm(heights$weight, mle_female$mean_weight, mle_female$sd_weight))
predict <- ifelse(num/den > 0.5, "Male", "Female")

# confusion matrix
table(predict, heights$sex)
```

```
##
## predict  Female Male
##   Female    190    50
##   Male      48   762
```

# Linear Discriminant Analysis (LDA)

- Naïve Bayes is pretty naïve
  - assumes independence of features
- LDA relaxes this assumption
- LDA still assumes that features are normally distributed
  - In particular, that they are **Multivariate normal**:  $\mathbf{x}_i | C = k \sim N(\boldsymbol{\mu}_k, \boldsymbol{\Sigma})$
  - $\boldsymbol{\Sigma}$  is the covariance matrix (defines relationship among features; in Naïve Bayes, can write as a diagonal matrix)
  - LDA assumes  $\boldsymbol{\Sigma}$  is the *same for each class*

# LDA classifier


- $p(C = k)$  and  $\mu_k$  estimated from MLE as in Naïve Bayes
- Pooled Sample covariance estimate  $\hat{\Sigma} = \frac{1}{n} \sum_{k=1}^K \sum_{i:c_i=k} (\mathbf{x}_i - \mu_k)(\mathbf{x}_i - \mu_k)^T$
- Predict class  $k$  for observation  $i$  that maximizes

$$p(C = k) f_{MVN}(\mathbf{x}_i | \hat{\mu}_k, \hat{\Sigma})$$

where  $f_{MVN}(\mathbf{x} | \mu, \Sigma)$  is the probability density function of the Multivariate Normal distribution

- Turns out that the decision rules for classifying into class  $k$  versus  $k'$  are **linear** combinations of the predictors
- If we instead let  $\Sigma$  be different for each class  $k$ , this is **Quadratic discriminant analysis**: decision boundaries are quadratic curves

# LDA vs QDA





# LDA in height/weight example


```
height_lda = lda(sex ~ weight + height,  
                 data = heights)  
height_lda
```

```
## Call:  
## lda(sex ~ weight + height, data = heights)  
##  
## Prior probabilities of groups:  
##      Female      Male  
## 0.2266667 0.7733333  
##  
## Group means:  
##           weight    height  
## Female 137.0573 64.93942  
## Male   174.2648 69.31475  
##  
## Coefficients of linear discriminants:  
##           LD1  
## weight  0.063123732  
## height -0.005069971
```

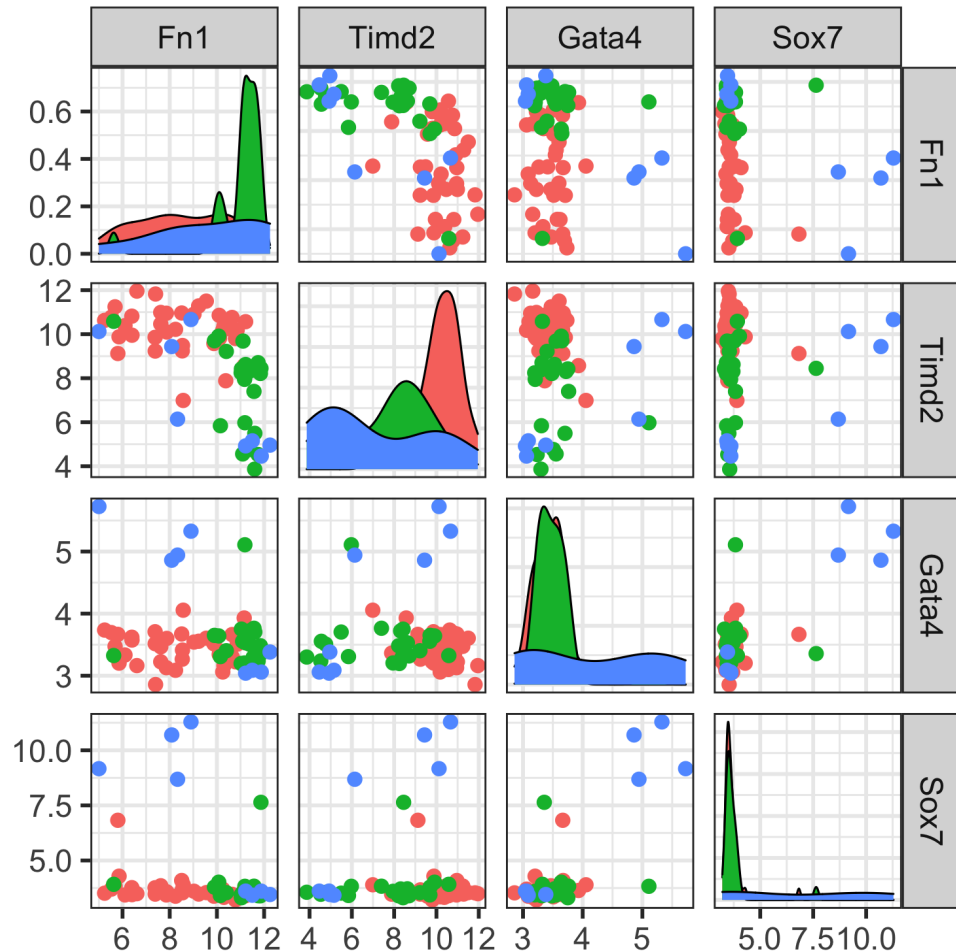
```
# confusion matrix  
table(predict(height_lda)$class,  
       heights$sex)
```

```
##  
##           Female Male  
## Female      186   47  
## Male        52  765
```

# Visualization of decision boundary



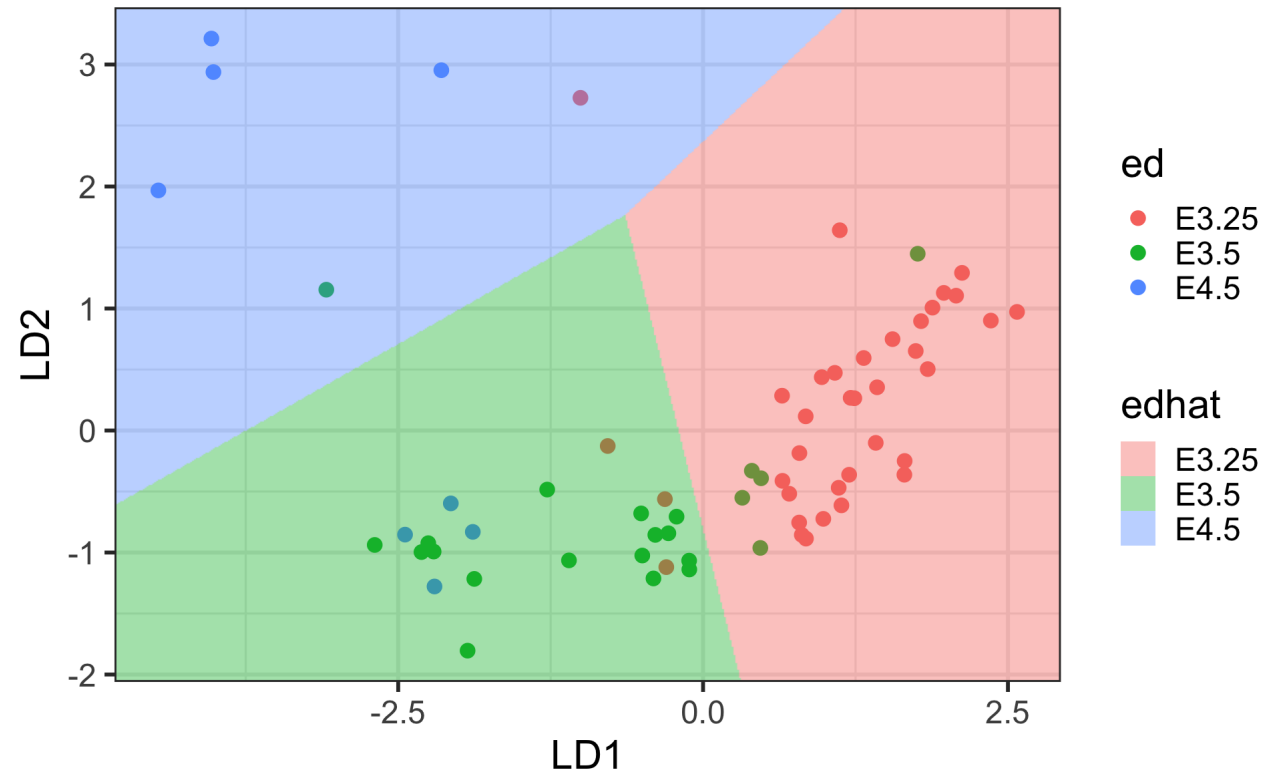
# Generalizing LDA to more than 2 classes



- We'd like to build a classifier that predicts the embryonic cell state from gene expression
- Specifically, we'd like to classify observations into one of three developmental time points (embryonic days: E3.25, E3.5, E4.5, shown in colour)
- Suppose that we already know that four particular genes (Fn1, Timd2, Gata4 and Sox7) are relevant to the task

Data/example source: **Modern Statistics for Modern Biology** by Holmes and Huber

# Generalizing LDA to more than 2 classes



# Logistic regression

- Logistic regression can be used as a binary (two-class) classifier
- Logistic regression vs two-class LDA
  - Both logistic regression and LDA use MLE to estimate  $p(C = k|\mathbf{x}_i)$
  - LDA assumes a parametric distribution for  $p(\mathbf{x}_i|C = k)$ ; if assumption is reasonable, can be more powerful
  - Logistic more resistant to outliers, model misspecifications


# Non-parametric

Why should we bother thinking about a model for  $p(C = k | \mathbf{x}_i)$  if our goal is just to partition the input space?

For example, if we want to assign a class to observation  $\mathbf{x}_i$ , why don't we just look at the class assigned to the point(s) **closest** to  $\mathbf{x}_i$ ?

# K-Nearest neighbour classifier

- One way to define "closest" is to use a fixed number of neighbours ( $K$ )
- Then count how many points of each class there are among these closest  $K$  neighbours to  $\mathbf{x}_i$
- Use the majority class as the predicted class



# Visualization of KNN classifier boundary




Fig 3b: Bzdok et al. (2018)



# Additional Resources

- Conceptual overview: [Chapter 12 of Modern Statistics for Modern Biology by Holmes and Huber](#)
- More detailed overview + R implementation: [Chapter 31 of Intro to Data Science by Irizarry](#)
- Mathematical framework: [Chapters 4 and 13 \(plus other chapters that expand on methods we didn't cover: 12 and 15\) in Elements of Statistical Learning by Hastie, Tibshirani and Friedman](#)